

YAKUP KORAY BUDANAZ

SoftHier Progress Report February 10



Overview of the Topics:

- **Programming Models**
- **DaCe + SoftHier**
- **Outlook For Next Weeks**

- **AI Accelerators and Programming Models**

AI Accelerators – An Overview of Hardware Features

Accelerator	Cache	Programmable NoC	Explicit Memory Levels	Dedicated MMU	Oversubscription to PEs	Double Precision Support
Nvidia A100	✓	×	2	✓	✓	✓
Nvidia H100	✓	✓*	2	✓	✓	✓
Ascend 910B	×	×	3	✓	×	×
Graphcore IPU Mk2	×	✓	2	×	✓	×
AMD VC2802**	×	✓	2	×	×	×
Microsoft Maia 100***	×	?	?	✓	?	×
Cerebras WSE-2	×	✓	2	×	×	×
Intel Gaudi v3	✓	×	2	✓	✓	×
AWS Trainium2	×	×	2	✓	×	×

Information collected from available white-papers, official tutorials and HotChips presentations.

AI Accelerators

- TPU is not there because I found nothing on the low-level programming of TPUs.
- Only see how to write in High Level Operator (HLO) IR for Accelerated Linear Algebra (XLA) DL compiler used as backend for JAX and TensorFlow.

Programming Models: High-Level IRs:

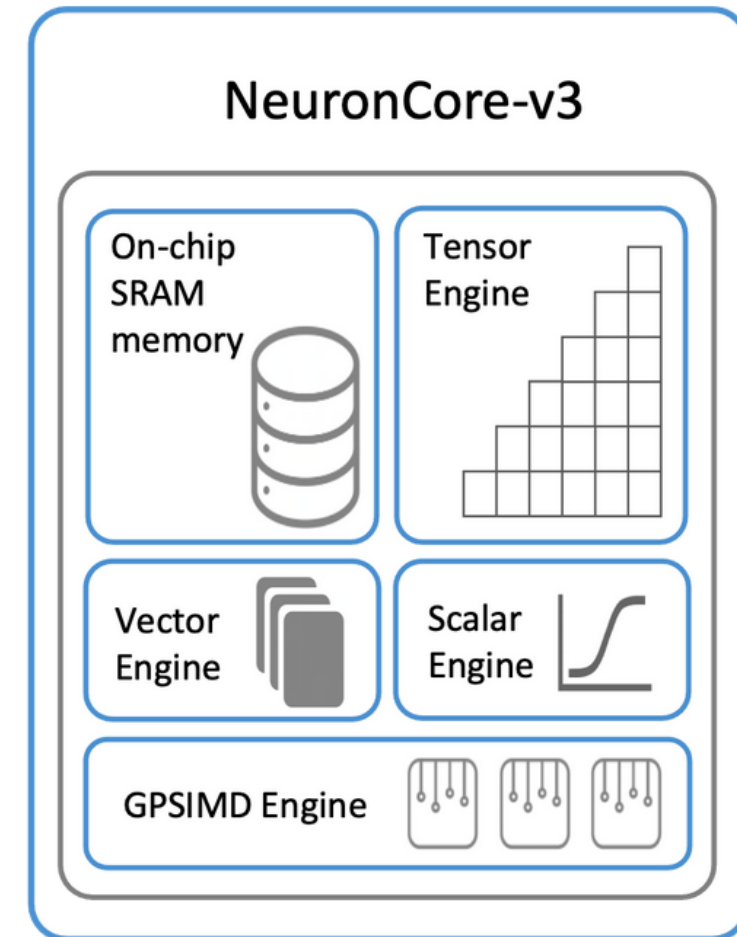
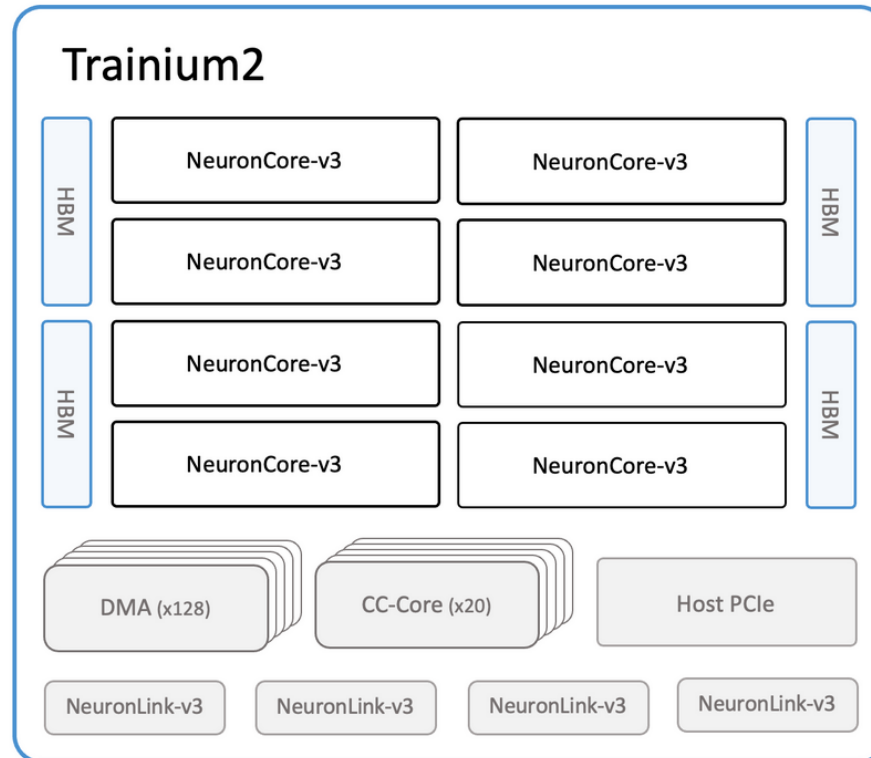
- A data-flow graph consisting of operations and input/outputs.
- Not really different a graph of consisting of ONNX operators. It is a variant of MLIR. Looks like this:

```
1  func @add_tensors(%arg0: tensor, %arg1: tensor) -> tensor {  
2      %0 = "mhlo.add"(%arg0, %arg1) : (tensor, tensor) -> tensor  
3      return %0 : tensor  
4  }
```

Abstracting Away Complexity – Method 1: Provide a list of Operations on Tensors + Efficient Implementations for them

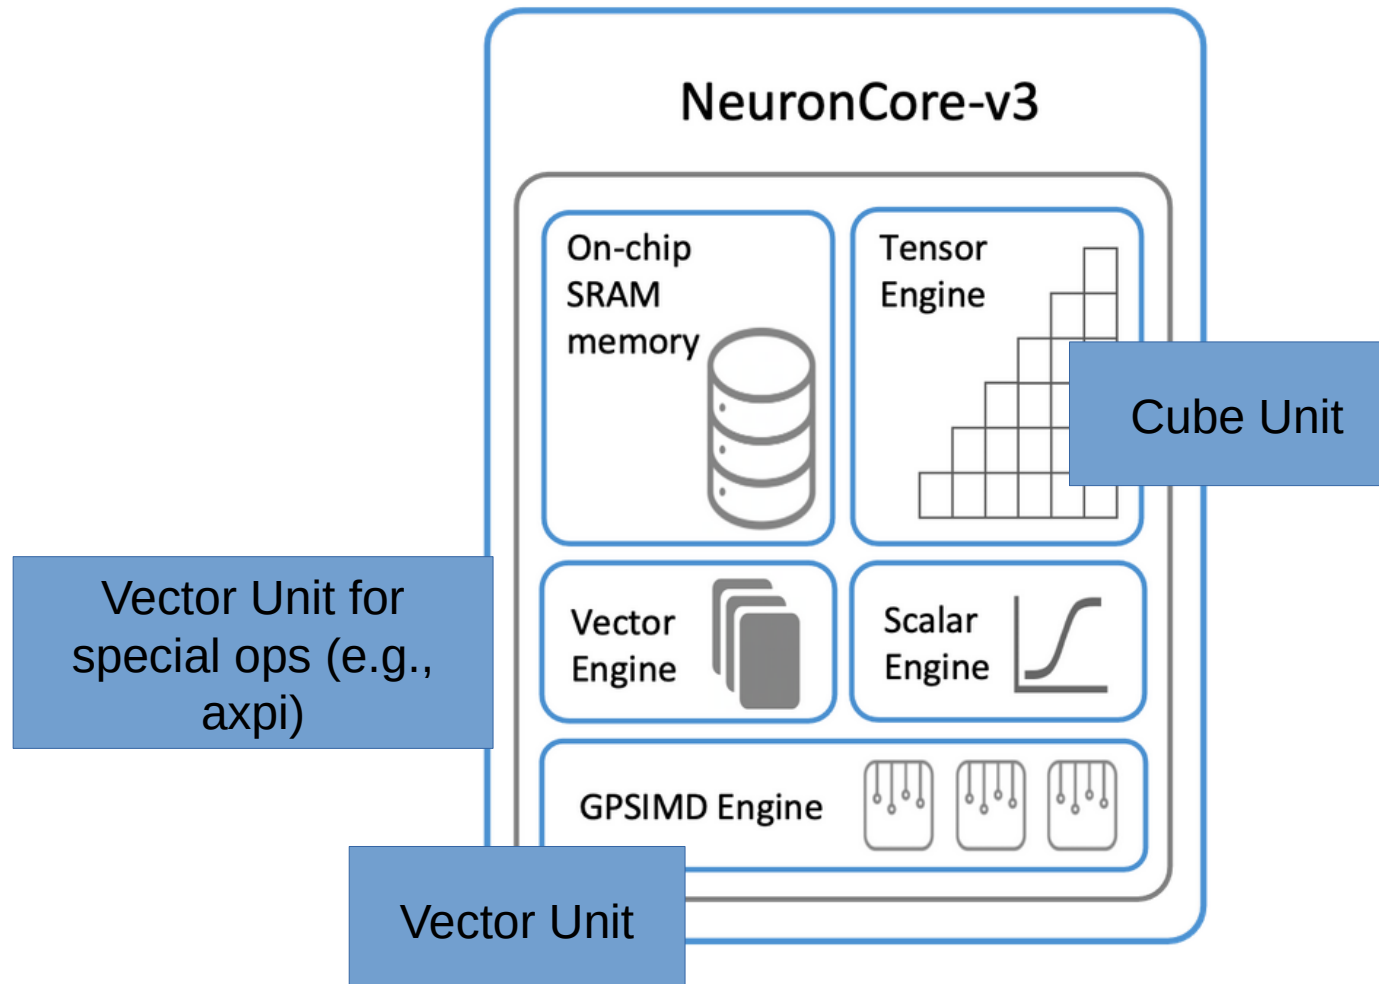
Google TPU only supports this, most AI accelerators support this

A very similar looking device - AWS Trainium2




<https://awsdocs-neuron.readthedocs-hosted.com/en/latest/general/arch/neuron-hardware/trainium2.html>

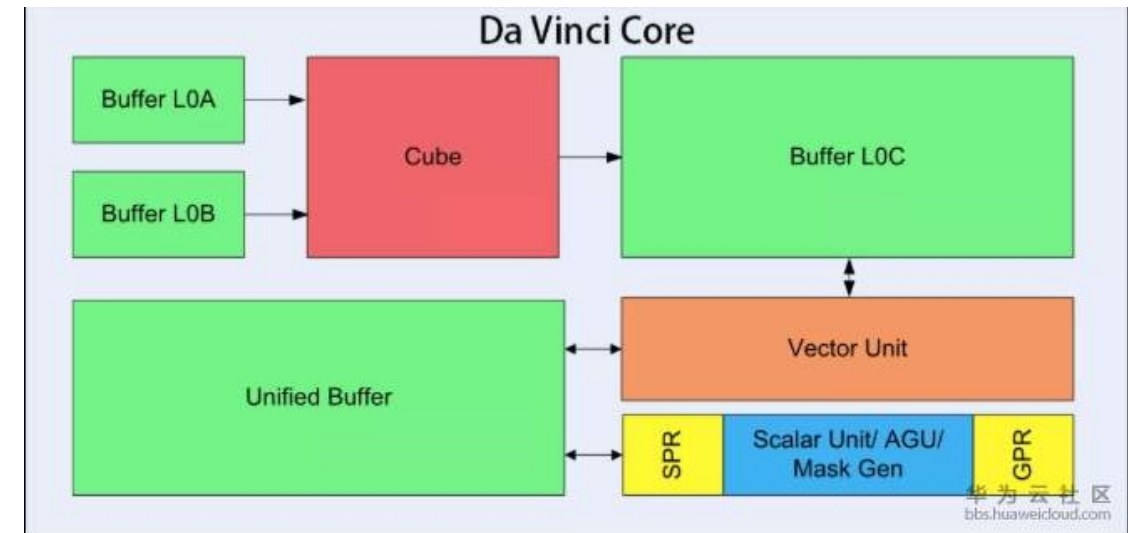
A very similar looking device - AWS Trainium2



A very similar looking device - AWS Trainium2



https://awsdocs-neuron.readthedocs-hosted.com/en/latest/general/nki-programming_model.html#nki-programming-model



<https://bbs.huaweicloud.com/blogs/106229>

A very similar looking device - AWS Trainium2

- Provides a tile-based programming interface similar to Triton named NKI
- No low-level programming interface for other computational units.

Abstracting Away Complexity – Method 1: Provide a list of Operations on Tensors + Efficient Implementations for them

Programming Models

- Betting on providing efficient implementations for specific tensor operations, or betting on tile-based languages.

Triton

- Load tiles and describe computation on tiles.
- Hopefully Triton back end + transformations will come-up with the most efficient way to move the memory

Abstracting Away Complexity – Method 2: Tile-Based language + a lot of memory movement optimizations

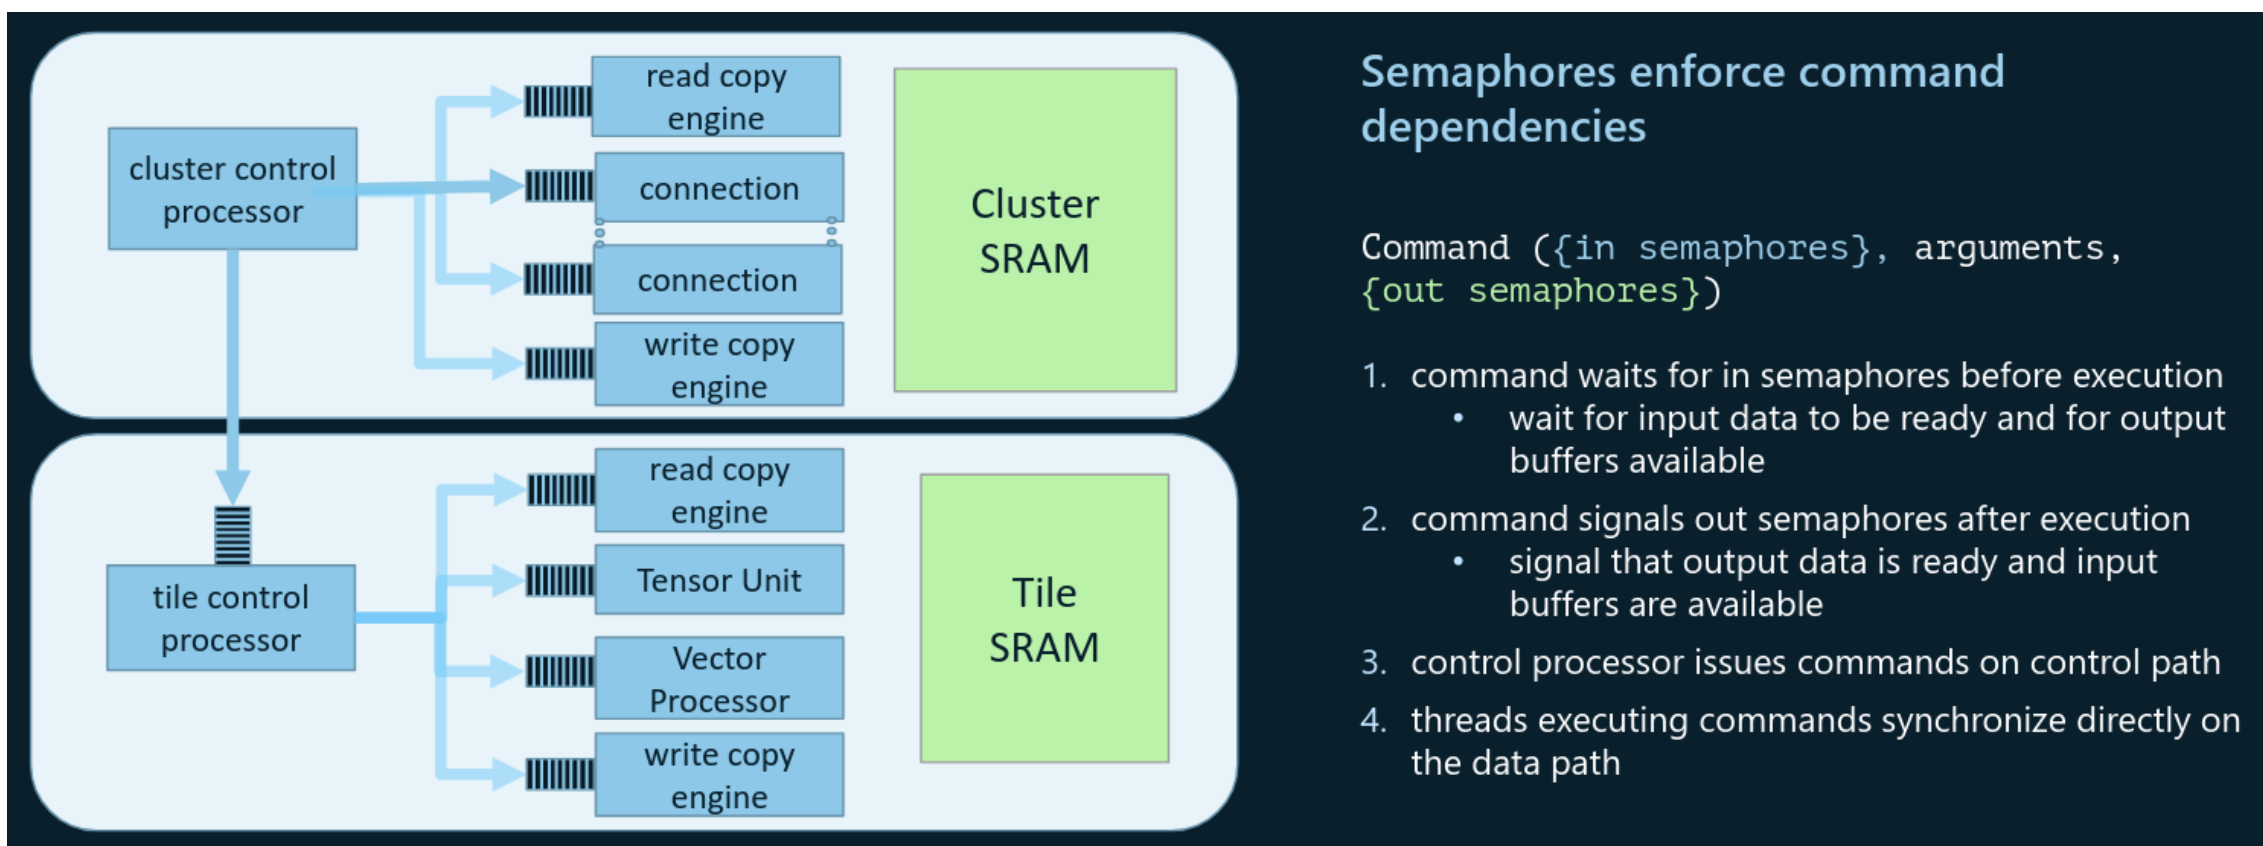
Triton

- Tiled-based CUDA Wrapper (+embedded DSL in Python to generate Triton IR and into LLVM IR to ptx)
- Example stencil kernel.

```
1  A_bottom = triton.language.load(  
2      A_ptr + (i + 2) * N + (j + 1), mask=(i < N - 2) & (j < N - 2), other=0.0  
3  )  
4  
5  ...  
6  B_new = 0.20 * (A_left + A_right + A_top + A_bottom + A_center)  
7  triton.language.store(  
8      B_ptr + (i + 1) * N + (j + 1), B_new, mask=(i < N - 2) & (j < N - 2)  
9  )  
10
```


Microsoft Maia 100

- From the hotchips presentation. Looks like Maia will have a model similar to asynchronous tasks, dependencies will be satisfied using semaphores.



Semaphores enforce command dependencies

Command ({in semaphores}, arguments, {out semaphores})

- command waits for in semaphores before execution
 - wait for input data to be ready and for output buffers available
- command signals out semaphores after execution
 - signal that output data is ready and input buffers are available
- control processor issues commands on control path
- threads executing commands synchronize directly on the data path

Microsoft Maia 100

- It also plans to provide Triton as a programming model

Abstracting Away Complexity – Method 3: Asynchronous Tasks Based Language?

Intel Gaudi 3

- It provides TPC language for low-level programming of the device
- Almost identical to CUDA with more intrinsics. It exposes warp (vector processor) more than the thread.

Abstracting Away Complexity – Method 4: Copy CUDA

Graphcore:

- Provide poplar C++ API to create a dataflow graph consisting of tensor operations and input-output dependencies.
 - Looks like SDFG with library-nodes only.
 - The operation implementation uses a BSP-like model, the communication is implicitly generated using tiles. (It describes how computation is mapped to cores using a tile abstraction)
 - Inter-PE communication is facilitated by the compiler (derived from tiling in form of graph partitioning).
 - For more explicit control they provide assembly for their ISA.
- The **bulk-synchronous parallel** model of execution. This decomposes the runtime execution into three phases: local compute, global synchronisation, and data exchange.
 - The graph representation of computations. The Poplar graph programming framework operates on a computational graph where vertices represent operations and edges represent the input and output data.

Graphcore

Abstracting Away Complexity – Method 5: Graph-based language + Graph-partitioning

AMD FPGAs

- Super low-level
- Super chaotic report
- Looks like assembly with some wrappers
- Inter-PE communication through an API similar to streams
- Provides a dataflow graph abstraction for operators

Abstracting Away Complexity – Method 6: Just Don't?

Asynchronous Memory Pipelines

- Pipeline as a FIFO Queue following producer-consumer pattern (asynchronous).

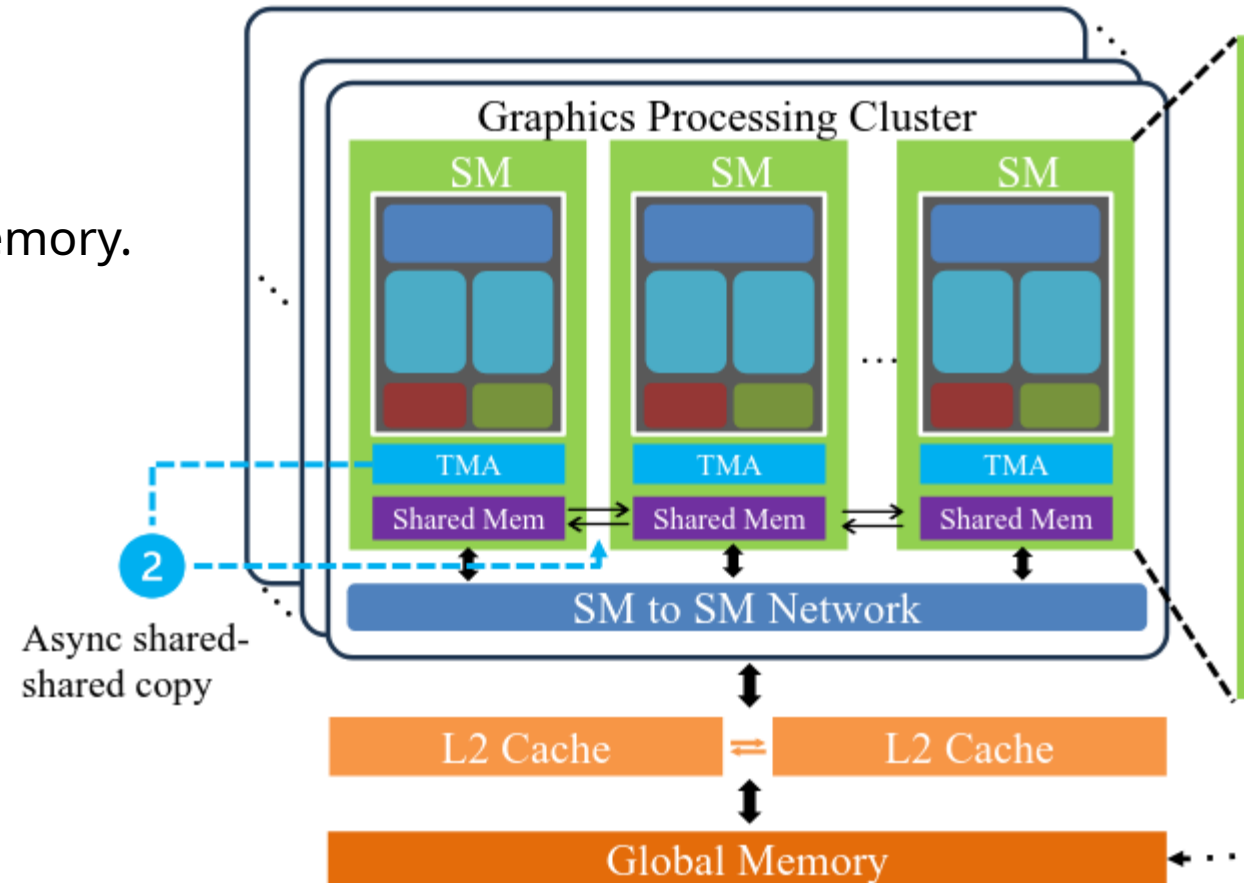
```

1  cuda::pipeline<cuda::thread_scope_thread> pipe = cuda::make_pipeline();
2
3  cuda::memcpy_async(&smem[stage][threadIdx.x], &src[idx], sizeof(int), pipe);
4      pipe.producer_commit();
5
6  for (...) {
7      cuda::pipeline_consumer_wait_prior<num_stages - 1>(pipe);
8      __syncthreads();
9
10     compute();
11
12     __syncthreads();
13     pipe.consumer_release();
14
15     pipe.producer_acquire();
16     cuda::memcpy_async(&smem[stage][threadIdx.x], &src[idx], sizeof(int), pipe);
17
18     pipe.producer_commit();
19
20     stage = (stage + 1) % num_stages;
21 }
22 }
```

Asynchronous DMA Engine:

- Tensor cores are now asynchronous too.
- One TMA per SM to asynchronously move memory.

Arch	Precision	Programmability	Mode
Ampere	FP16,BF16, TF32,FP64, INT8,INT4,Binary	C: wmma PTX: mma, mma.sp	Sync
Ada	FP16,BF16,FP8, TF32,FP64, INT8,INT4,Binary	C: wmma PTX:mma, mma.sp	Sync
Hopper	FP16,BF16,FP8,TF32, FP64,INT8,Binary	C: wmma PTX: mma, mma.sp PTX: wgmma, wgmma.sp	Sync ASync



Asynchronous DMA Engine:

- CuTe / CUTLASS provides convenient ptx wrappers for common optimizations such as memory swizzling, and calls to tensor cores / TMA.

Abstracting Away Complexity – Method 7: Providing Assembly Wrappers for Common Use-cases

<https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html?highlight=TMA#using-tma-to-transfer-one-dimensional-arrays>

Asynchronous DMA Engine:

- TMA without high-level wrappers and only with ptx wrappers:
- Create, and copy tensor map
- Create many barriers, semaphores
- Async copies, set barriers, wait on barriers
- ...

Abstracting Away Complexity – Method 7: Even more convenient(?) assembly wrappers

<https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html?highlight=TMA#using-tma-to-transfer-one-dimensional-arrays>

My Take on Programming Models:

- If I were to design a programming language I would only treat asynchronous tasks (defined over tiles) with input and output dependencies and construct a graph of asynchronous tasks.
- Graph IRs are technically asynchronous. The dependencies are expressed as edges. The compiler or transformations need to find optimal memory movement pattern.

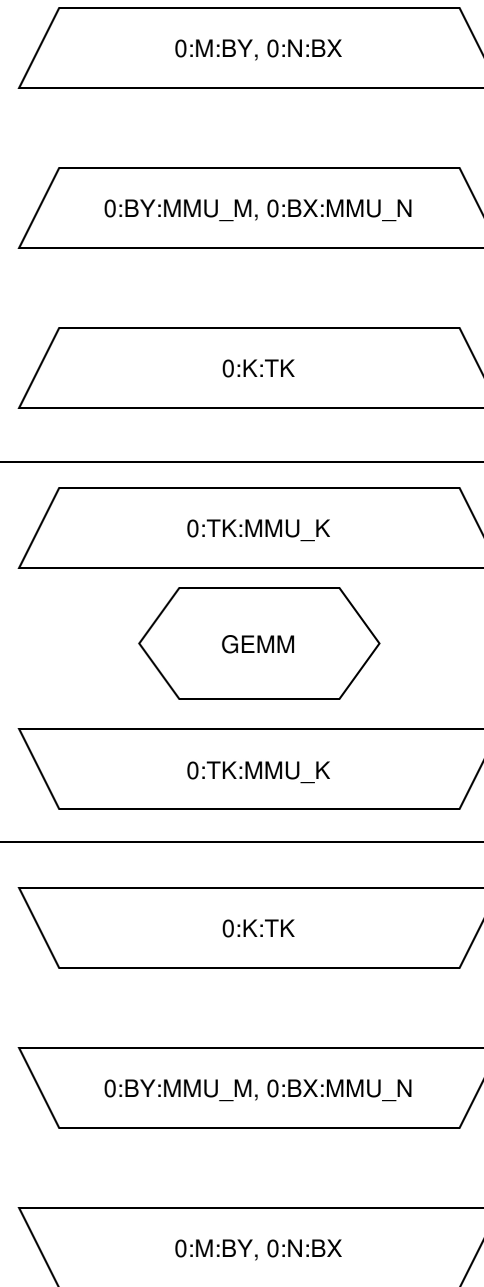
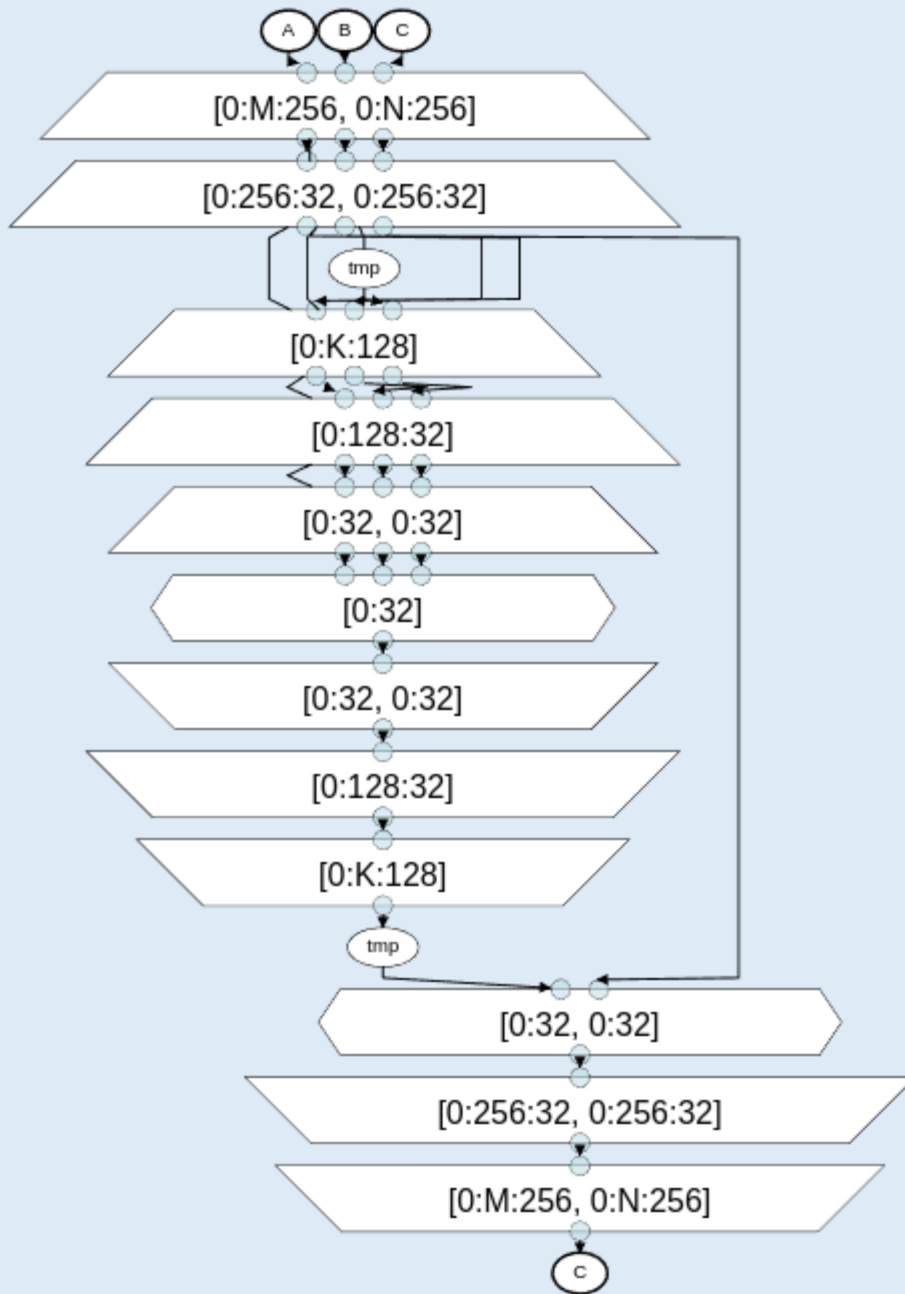
- **DaCe + SoftHier**

Expressing NoCs in SDFGs:

- In my opinion, the best way of expressing NoC communication is using streams.

Expressing NoCs in SDFGs:

- In my opinion, the best way of expressing NoC communication is using streams.
- Stream abstraction used in SDFGs is a FIFO queue with multiple senders and one receiver.
- How to transform SDFGs?
 - Designed a transformation that transforms a subgraph to use streams using a BSP-model-like description

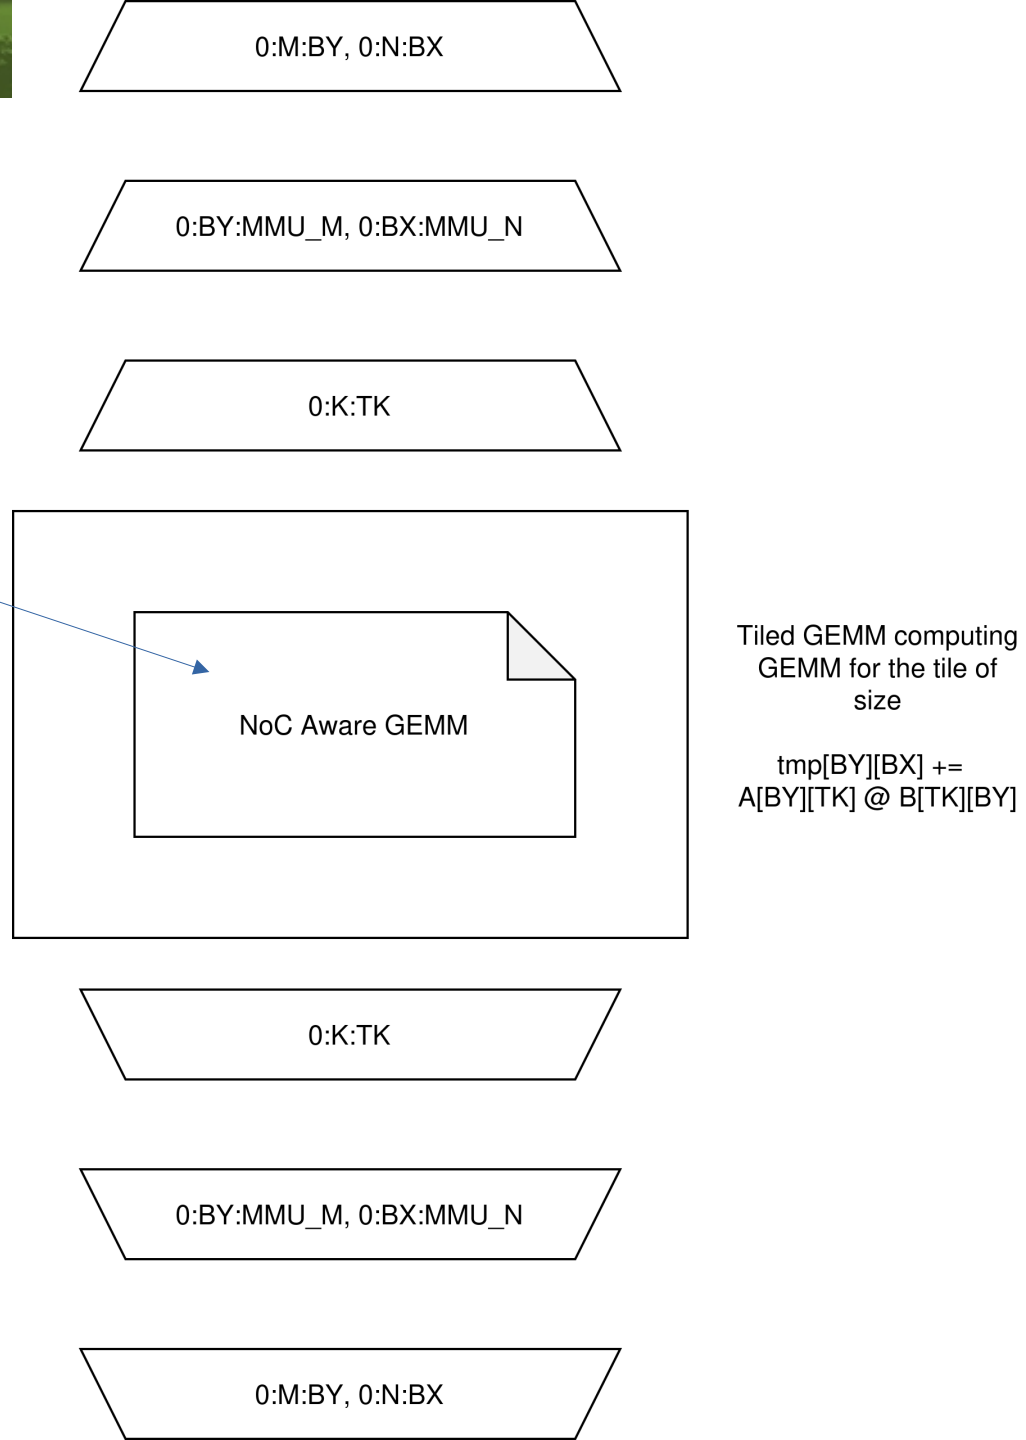


Tiled GEMM computing
GEMM for the tile of
size

$tmp[BY][BX] += A[BY][TK] @ B[TK][BY]$

Input is a BSP-model based description of the algorithm running inside.

- 1. Initial data dist (which core owns data)
- 2. Communication scheme between cores
- 3. Synchronization (if necessary in the architecture)



1. Initial distribution:

sA and sB are 2 dimensions streams, both of size $[N][N]$,
 $N = N_X = N_Y$

$A[j * \text{tileSize} + ((i+j) \% N) * \text{tileSize}] \rightarrow sA[i][((i+j)\%N)][:][:]$
 $B[((i+j) \% N) * \text{tileSize} + j * \text{tileSize}] \rightarrow sB[(i+j)\%N][j][:][:]$

2. Compute:

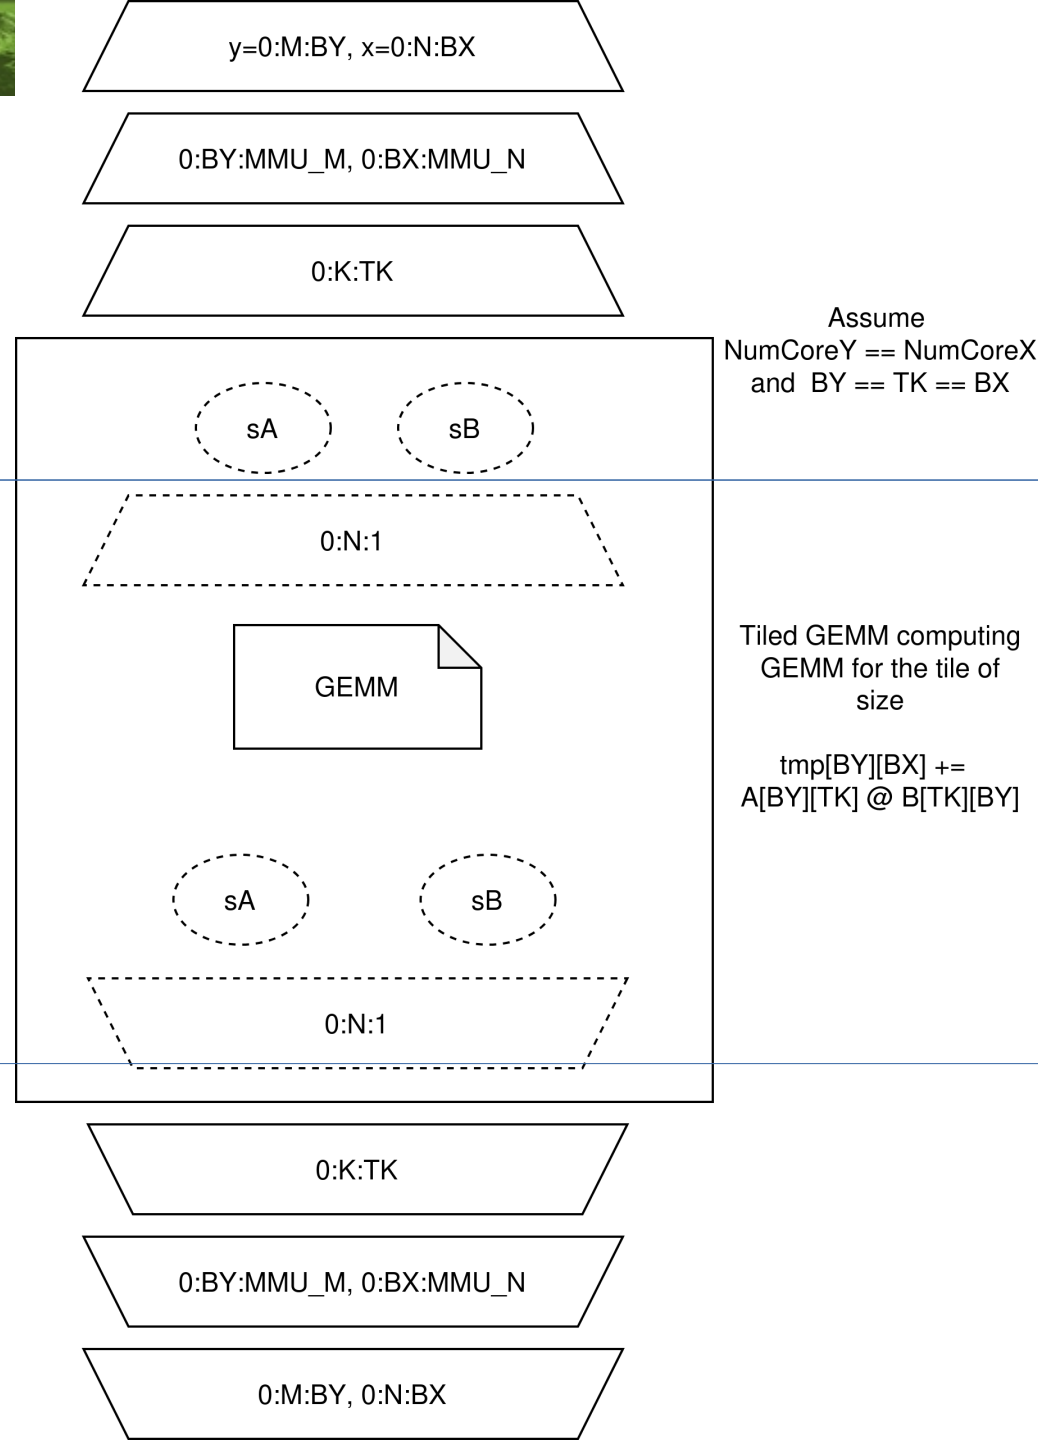
$\text{localA} = sA.\text{pop}()$, $\text{localB} = sB.\text{pop}()$
 $\text{tmp} += \text{localA} * \text{localB}$

3. Communication

$\text{LocalA}[:, :] \rightarrow sA[i][(j+N-1)\%N][:][:]$
 $\text{LocalB}[:, :] \rightarrow sB[(i+N-1)\%N][j][:][:]$

4. Synchronization

Yes/No depending on the architecture



No issues in part 1

2. Compute:

$\text{localA} = \text{sA.pop}()$, $\text{localB} = \text{sB.pop}()$
 $\text{tmp} += \text{localA} * \text{localB}$

3. Communication

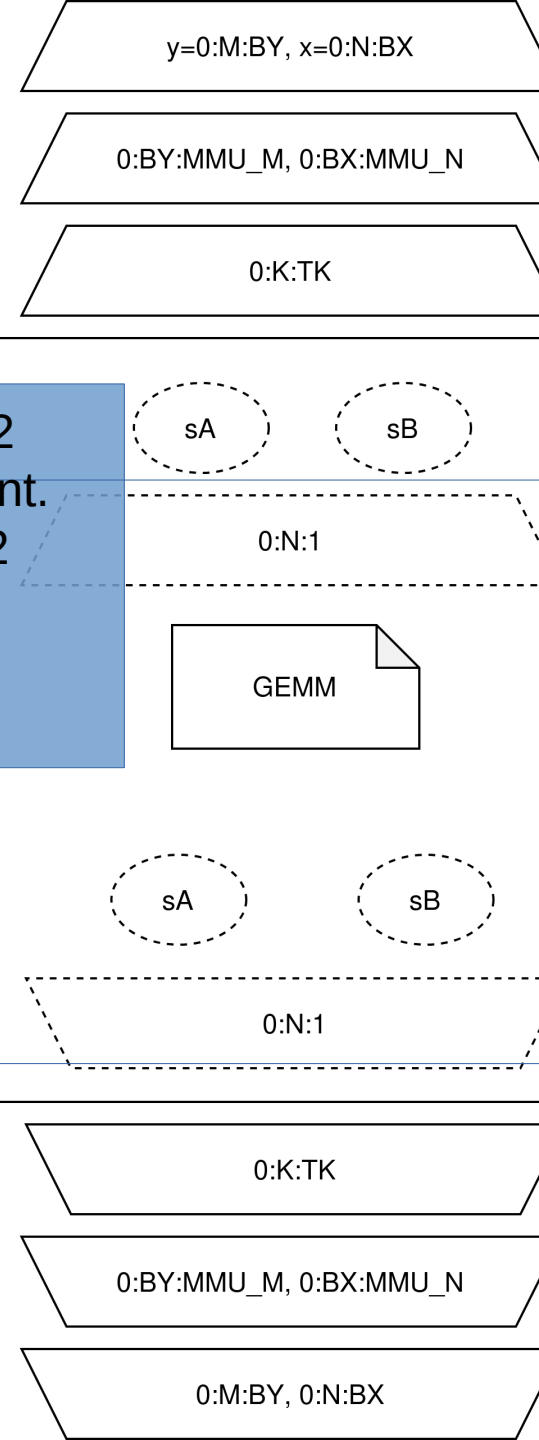
$\text{LocalA}[:, :] \rightarrow \text{sA}[i][(j+N-1)\%N][:][:]$

$\text{LocalB}[:, :] \rightarrow \text{sB}[(i+N-1)\%N][j][:][:]$

4. Synchronization

Yes|No depending on the architecture

For GEMM, the steps 2
 and 3 can be concurrent.
 For Stencil, the steps 2
 and 3 can't be
 concurrent. The user



Assume
 $\text{NumCoreY} == \text{NumCoreX}$
 and $\text{BY} == \text{TK} == \text{BX}$

Tiled GEMM computing
 GEMM for the tile of
 size

$\text{tmp}[\text{BY}][\text{BX}] +=$
 $\text{A}[\text{BY}][\text{TK}] @ \text{B}[\text{TK}][\text{BY}]$

No issues in part 1

2. Compute:
 localA = sA.pop(), localB = sB.pop()
 tmp += localA * localB

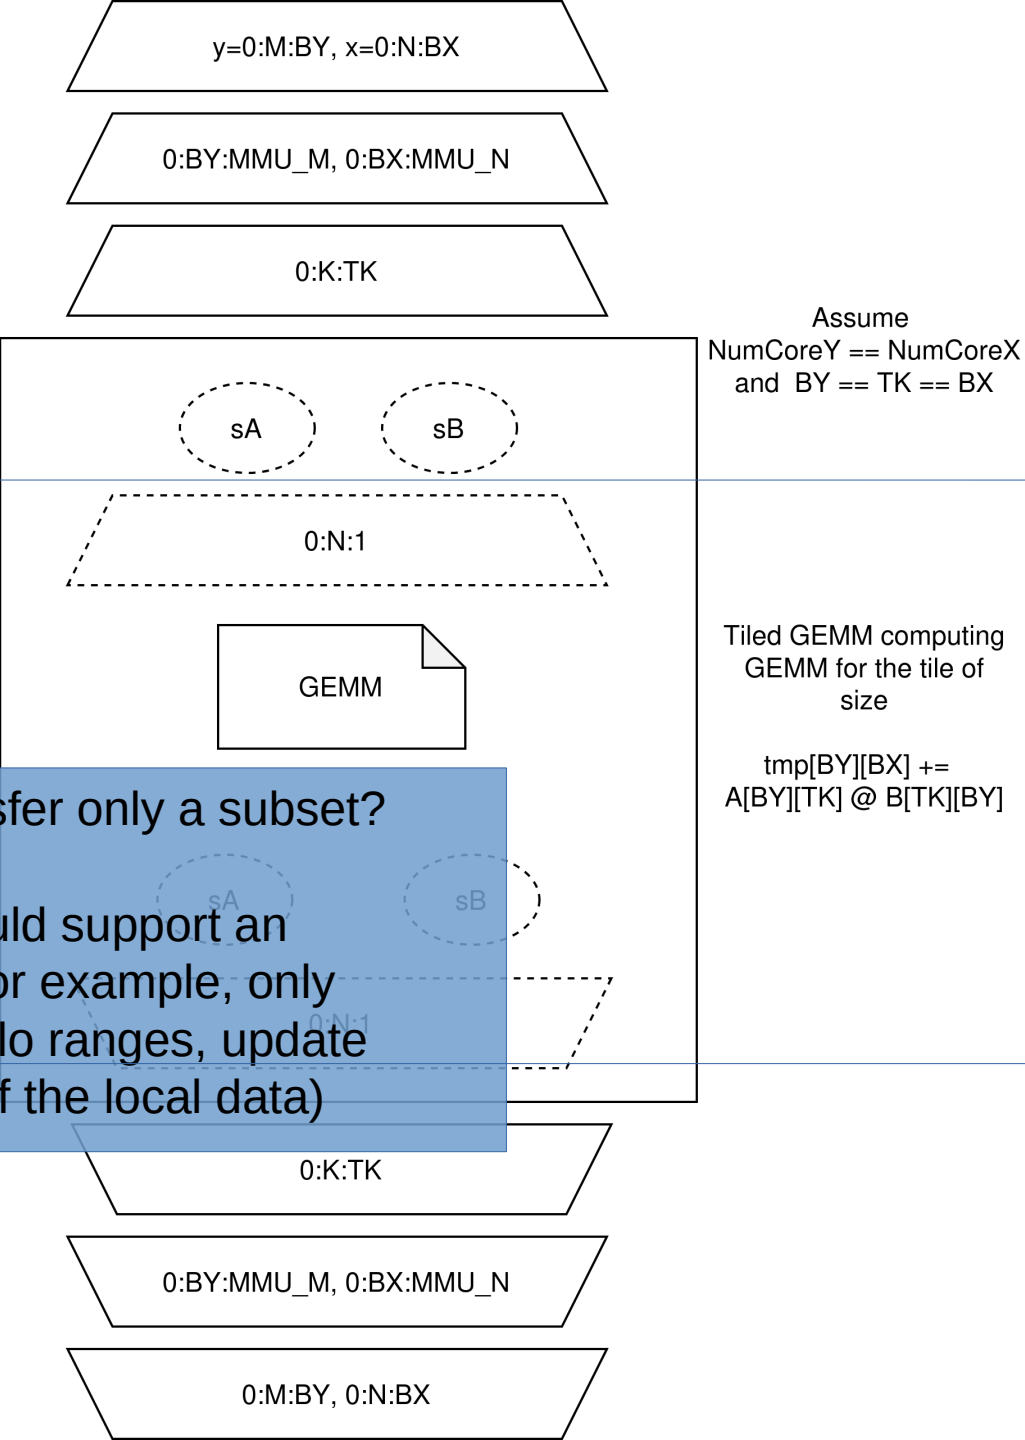
3. Communication

LocalA[:, :] → sA[i][(j+N-1)%N][:][:]
 LocalB[:, :] → sB[(i+N-1)%N][j][:][:]

No issues in part 4

What if need to transfer only a subset?

The description should support an update. (In stencil, for example, only communicate the halo ranges, update the local boundary of the local data)



1. Initial distribution:
 sA and sB are 2 dimensions streams, both of size [N][N], $N = N_x = N_y$,
 $A[i * tileSize][((i+j) \% N) * tileSize] \rightarrow sA[i][(i+j)\%N][:][:]$
 $B[((i+j) \% N) * tileSize][j * tileSize] \rightarrow sB[(i+j)\%N][j][:][:]$

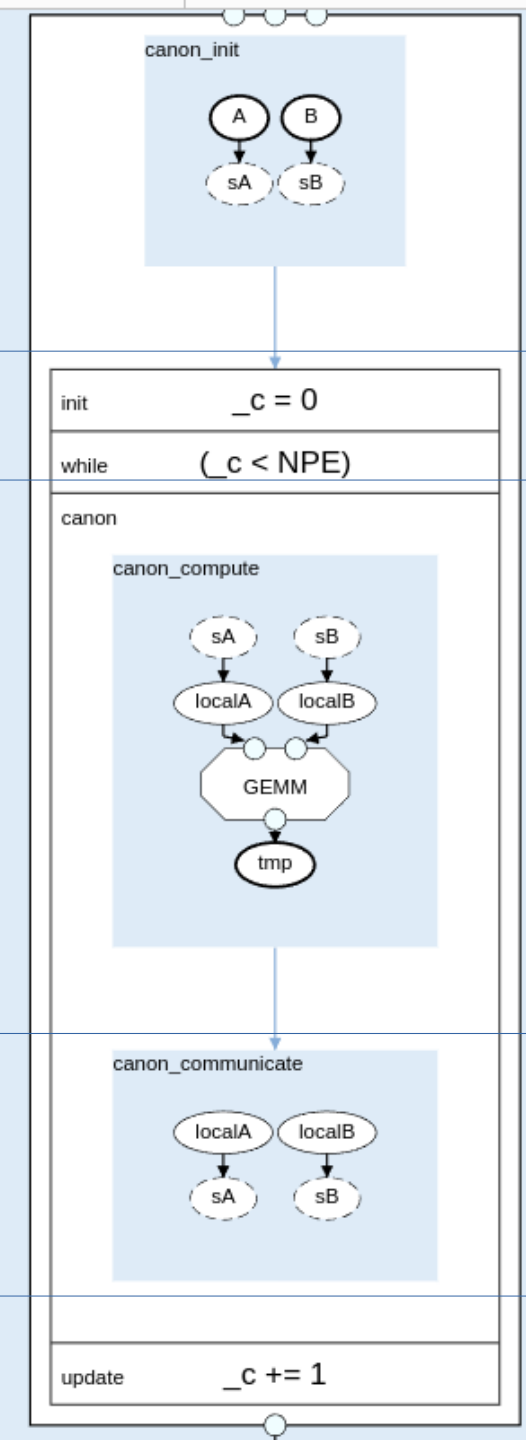
2. Loop Range

2. Compute:
 localA = sA.pop(), localB = sB.pop()
 tmp += localA * localB

3. Communication

$LocalA[:, :] \rightarrow sA[i][(j+N-1)\%N][:][:]$
 $LocalB[:, :] \rightarrow sB[(i+N-1)\%N][j][:][:]$

4. Synchronization
 Generated as a library node on a 3rd state or as a loop property



- Scalar → MMU Pass (Pattern detection + Reading Tasklet ASTs), working for couple of patterns
- Scalar → Vector Unit Pass (Pattern detection + Reading Tasklet ASTs), works for couple of patterns, need to extend patterns later as necessary
- GEMM Input Type information is added to access nodes which is used by the controller to apply the correct ExplicitMemoryMove transformation
- Passing “input purpose” to SDFGs using access nodes manually (done)
- Automating this information by extending Scalar → MMU Pass (mostly done)
- Pass to insert memory movement between different computational units (under work)

Outlook / TODOs:

- Work further on transformations
- Extend cube unit parts of the code-gen

My Take on Programming Models:

Sketched how it could look like. Many considerations are necessary. Message buffers, supporting out-of-order messages, low-level implementation of completions of one-sided communication.

```

1  __kernel__ vadd<N>(in A[N*1024], in B[N*1024], inout C[N*1024]){
2      decl id = PEs<N>
3      decl TS = 1024
4      a1 = reserve(localB[1024], TCM)
5      a2 = reserve(localA[1024], TCM)
6      a3 = reserve(localC[1024], TCM)
7      l1 = a1.then(load(in:A[i*TS:(i+1)*TS], out:localA[:]))
8      l2 = a2.then(load(in:B[i*TS:(i+1)*TS], out:localB[:]))
9      t1 = join(l1, l2).then(add(in:localA[:], in:localB[:], out:localC[:]))
10     s1, r1 = t1.then(send(to:(i+N-1)%N, in:localC[:], out:localC[:]))
11     c1 = r1.then(store(in:localC[:], out:C[i*TS:(i+1)*TS]))
12     return c1
13 }
14
15 ...
16
17 vadd_kernel = vadd<N>(A, B, C)
18 vadd_kernel.wait()

```