

YAKUP KORAY BUDANAZ

SoftHier May 19 Status Update



Overview of the Topics:

- **Fixes for AscendC Build Pipeline And Vector Unit**
- **New Strategy For Generating Copies**
- **Current Status of CubeUnit Codegen and Current Issues**

Work on the AscendC Backend:

- For performance numbers I refer to:
<https://cset.georgetown.edu/publication/pushing-the-limits-huaweis-ai-chip-tests-u-s-export-controls/>
- I see HBM capacity is 32 GB on the server there I assume, the server 910B has 800 GB/s HBM bandwidth.


	Ascend 910 series (first-generation)	Ascend 910B series (second-generation)
Launch Date	Mid 2019	Late 2022
Performance (FP16 TFLOPS)	220 - 320	280 - 400
Total Number of AI Cores	32	?
Number of Active AI Cores	30 - 32	20 - 25
Clock Speed	0.9 - 1.15 GHz	1.65 - 1.85 GHz
Maximum On-chip Memory	76 MB	211 MB
HBM Type	HBM2	HBM2e
HBM Bandwidth	1228 GB/s	800/1600 GB/s
HBM Capacity	32 GB	32/64 GB
Fabrication	TSMC 7nm (N7+)	SMIC 7nm (N+2)

Work on the AscendC Backend:

- Updated NPU-Info Parsing Script to work on 910B server.
- Updated CMake Build Files, Compile Flags etc.
- Fixed couple of compile errors that appeared after migrating to 8.1RC1 related to typecasts. (From 8.0)

Vector Unit Also Functional On Ascend 910B:

Achieved Bandwidth vs Vector Size for Each Fragment Size



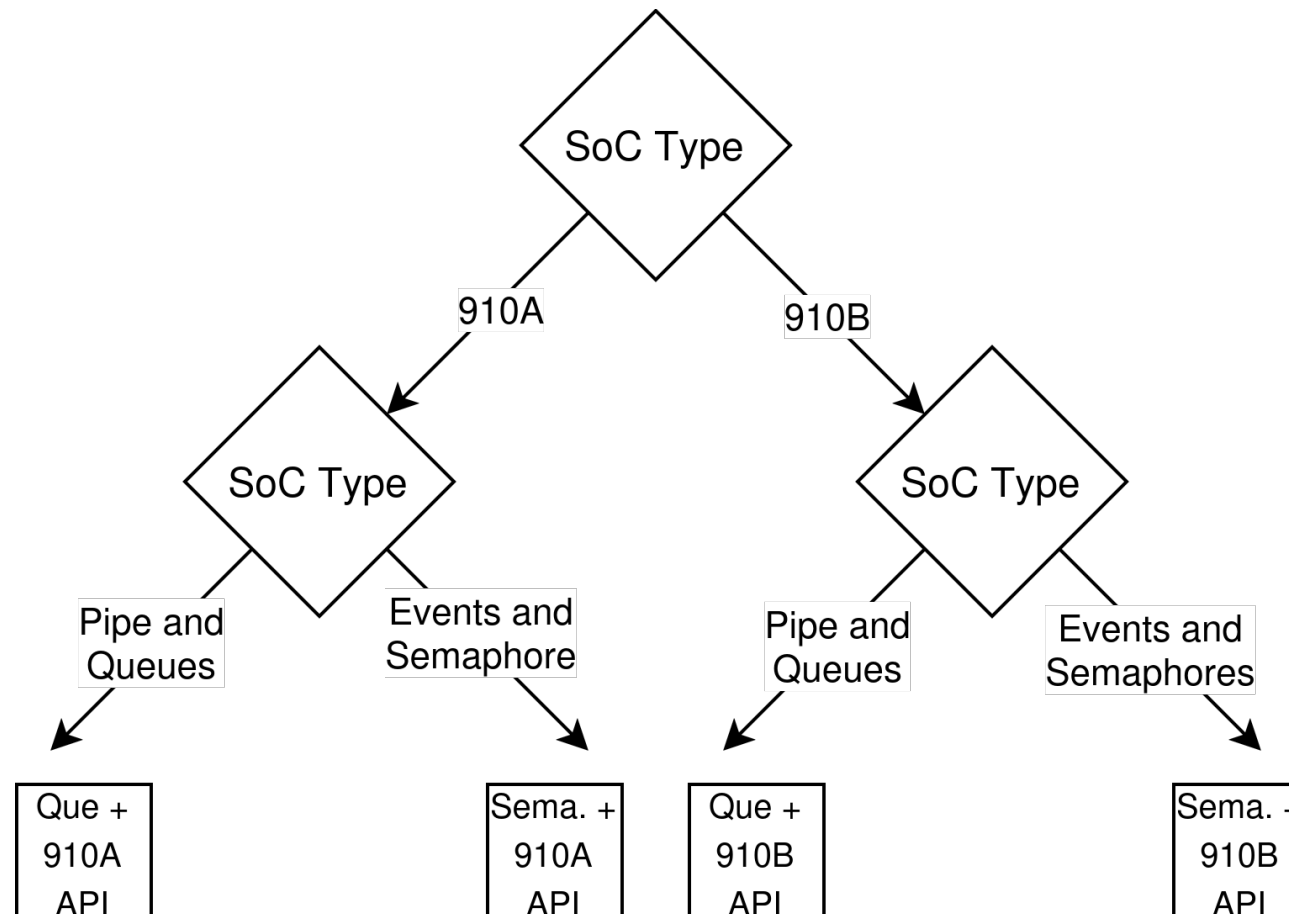
Vector Copy / Addition Benchmark

- Ran N=25+5 benchmarks, Outliers (Z-score > 3) removed, 100 bins, launched 2 kernels per program, measured the second kernel. Discarded the first 5 programs launches.
- The runtime distribution is not normal distribution. The median runtime of the N'=25 runs. The distribution looks binomial and the difference between minimum and maximum runtime is not significant.
- Vector Copy kernel computing $\mathbf{B} = \mathbf{A}$ timeouts using AscendC::Que's, the kernel computes $\mathbf{C} = \mathbf{A} + \mathbf{B}$
- Every AiCore adds **frag_size** elements at a time, vectors added $32 * \mathbf{frag_size}$ at a time. Pseudocode for the copy benchmark:

```
1  for (int64_t i = 0; i < vector_size; i += 32 * frag_size)
2  {
3      int ii = (frag_size * AscendC::GetBlockIdx());
4      {
5          ...
6          Add(OUT_frag_C, IN_frag_A, IN_frag_B, frag_size);
7          ...
8      }
9  }
```


New Strategy To Generate Copies:

- Strategy pattern with a pool of copy functions.
- One the copies are chosen depending on the configuration and the hardware model.




New Strategy To Generate Copies:

- Why a pool of copy functions?




Possible memory copies for 910A



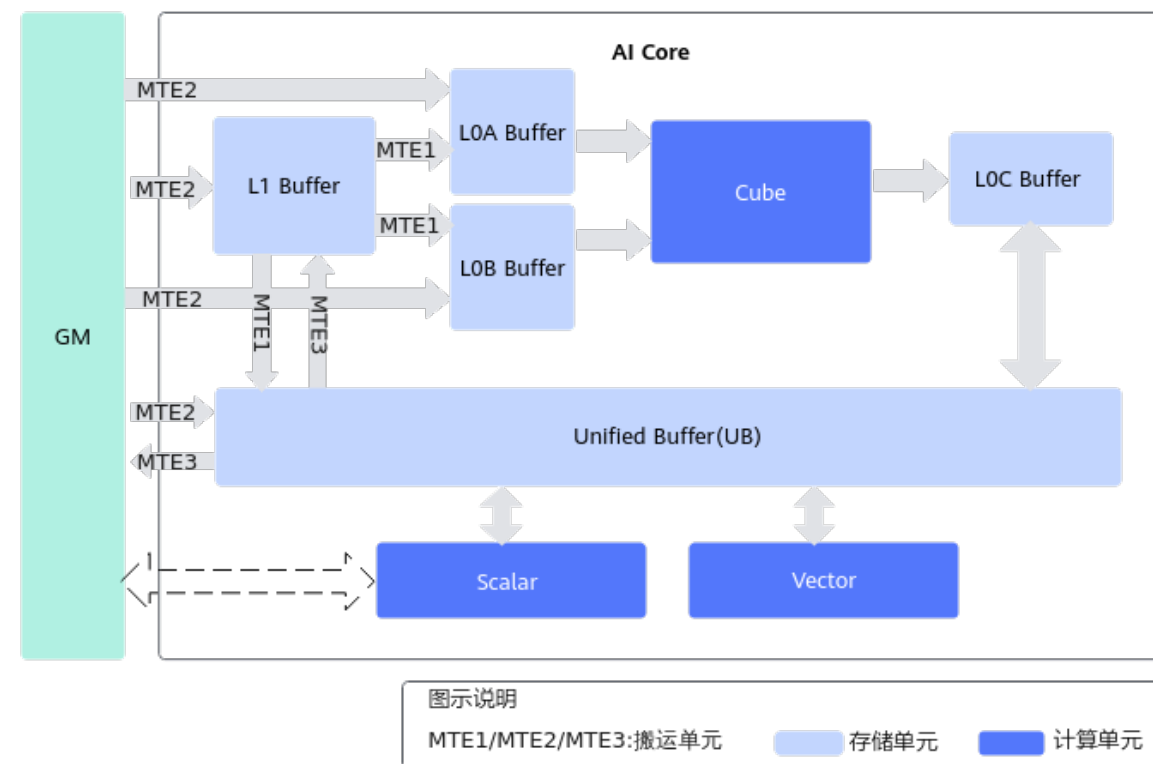
Possible memory copies for 910B

New Strategy To Generate Copies:

- Why a pool of copy functions?




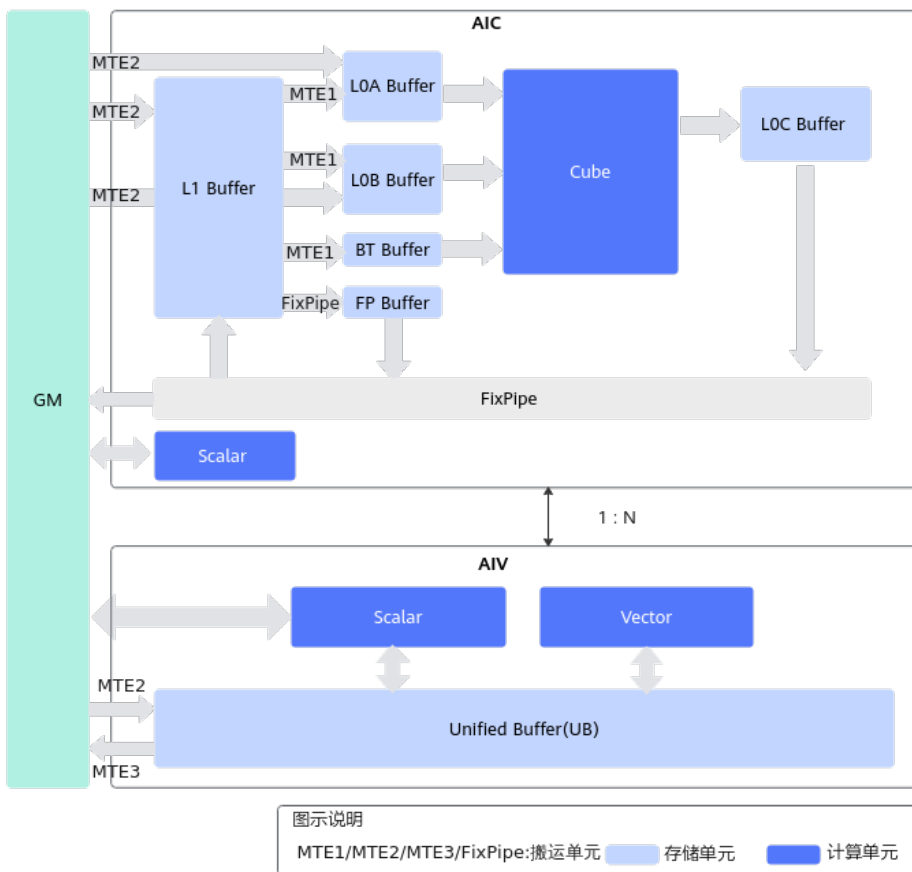
Possible memory copies for 910A



https://www.hiascend.com/document/detail/z/h/canncommercial/81RC1/developmentguide/opdevg/Ascendcopdevg/atlas_ascendc_10_0008.html

New Strategy To Generate Copies:


- Why a pool of copy functions?



Possible memory copies for 910B


New Strategy To Generate Copies:

- Why a pool of copy functions?



Possible memory copies for 910A

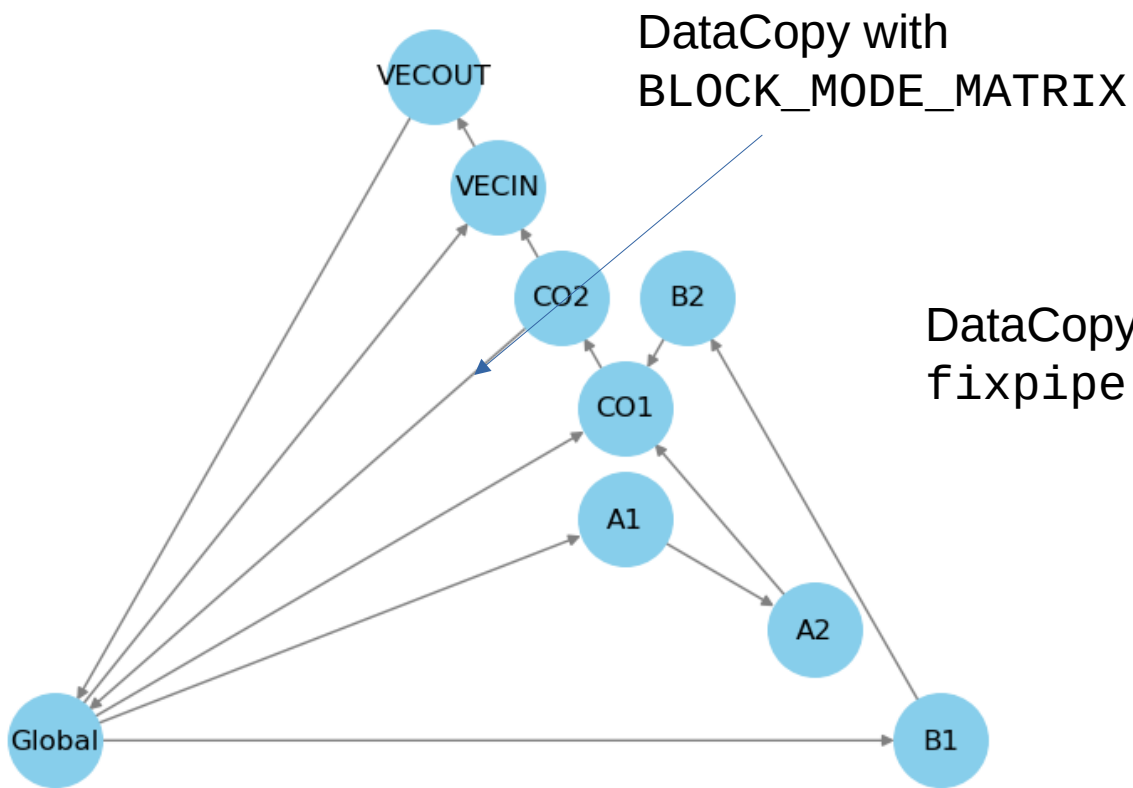
DataCopy with
For-loop + cache
lines width



Possible memory copies for 910B


New Strategy To Generate Copies:

- Why a pool of copy functions?



Possible memory copies for 910A

DataCopy with fixpipe



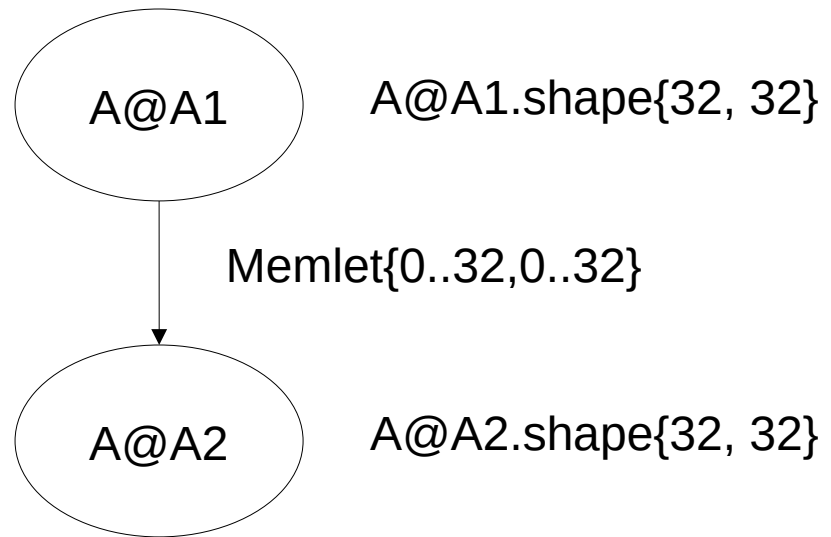
Possible memory copies for 910B

Motivation In the Backend Implementation:

- Expose the dataflow of data
 - While hiding the implementation specific layouts requirements from the user


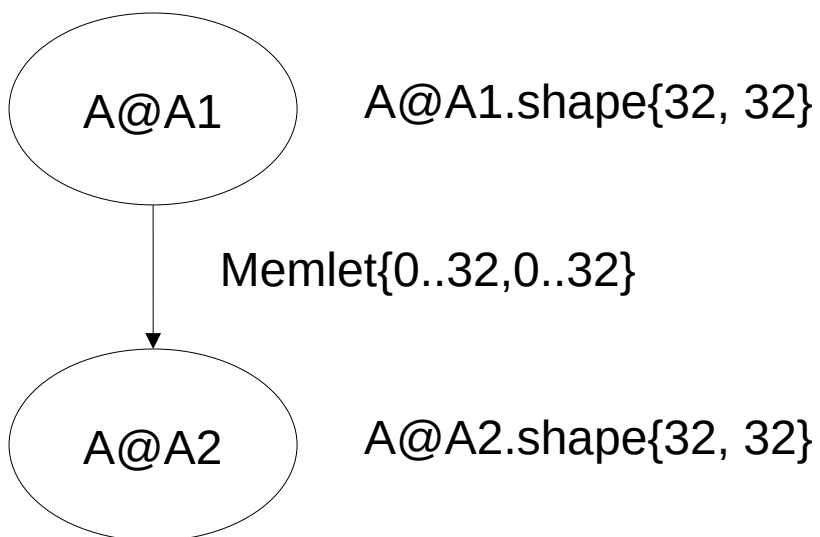
Anatomy of a Copy in SDFG:

The Goal is the show this to the user:



Anatomy of a Copy in SDFG:

The Goal is the show this to the user:



Such that the user does not need to worry about the which storage format is necessary for which A2 vs A1.

https://www.hiascend.com/document/detail/zh/canncommercial/81RC1/developmentguide/opdevg/Ascendcopdevg/atlas_ascendc_10_00006.html

Anatomy of a Copy in SDFG:

1. Alloc local tensor

LocalTensor<>(A_A1)


setGlobalBuffer(A, 32*32)

```

DataCopyParams Params;
Params.blockCount = 32;
Params.blockLen = 32 / 16;
Params.srcStride =
static_cast<uint16_t>((N / 16) - 1);
Params.dstStride = 0;
    
```

```

DataCopy(...)
QueueA_A1.Enqueue(A_A1)
    
```



Anatomy of a Copy in SDFG:


1. Alloc local tensor

LocalTensor<>(A_A1)

setGlobalBuffer(A, 32*32)

```
DataCopyParams Params;
Params.blockCount = 32;
Params.blockLen = 32 / 16;
Params.srcStride =
static_cast<uint16_t>((N / 16) - 1);
Params.dstStride = 0;
```

2. Read the
buffer size
read from the
memlet



```
DataCopy(...)
QueueA_A1.Enqueue(A_A1)
```

Anatomy of a Copy in SDFG:

1. Alloc local tensor

LocalTensor<>(A_A1)

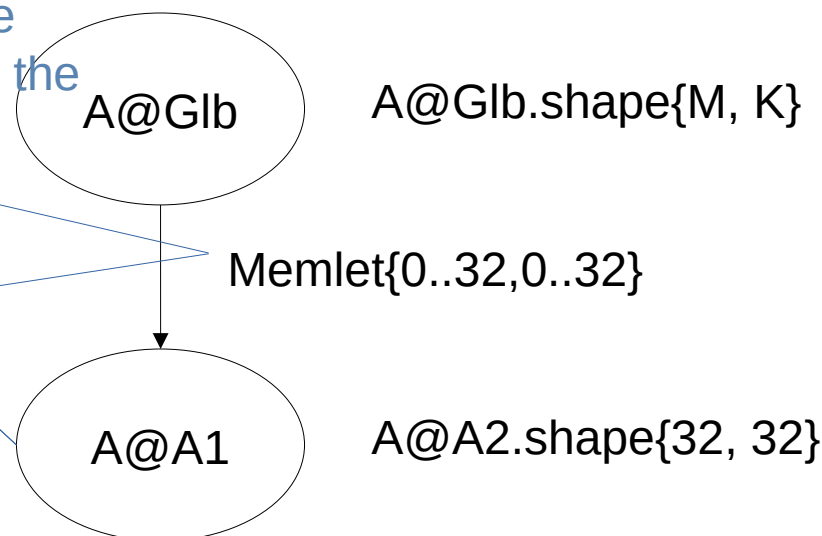
setGlobalBuffer(A, 32*32)

```
DataCopyParams Params;
Params.blockCount = 32;
Params.blockLen = 32 / 16;
Params.srcStride =
static_cast<uint16_t>((N / 16) - 1);
Params.dstStride = 0;
```

3. Read the copy parameters
(currently the input matrices
need to be row-major)

```
DataCopy(...)
QueueA_A1.Enqueue(A_A1)
```

2. Read the
buffer size
read from the
memlet



Anatomy of a Copy in SDFG:

1. Alloc local tensor

LocalTensor<>(A_A1)

setGlobalBuffer(A, 32*32)

```

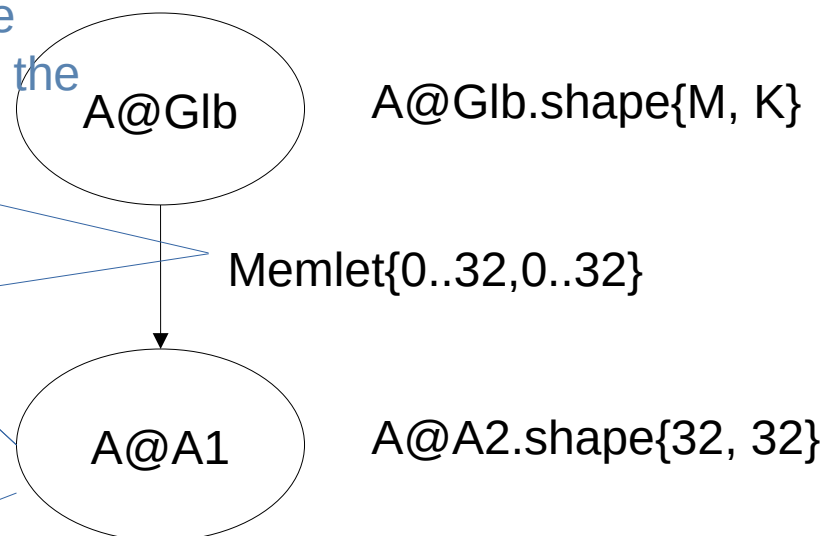
DataCopyParams Params;
Params.blockCount = 32;
Params.blockLen = 32 / 16;
Params.srcStride =
static_cast<uint16_t>((N / 16) - 1);
Params.dstStride = 0;
  
```

3. Read the copy parameters
(currently the input matrices
need to be row-major)

DataCopy(...)

QueueA_A1.Enqueue(A_A1) 4. Register copy and enqueue

2. Read the
buffer size
read from the
memlet



Anatomy of a Copy in SDFG:

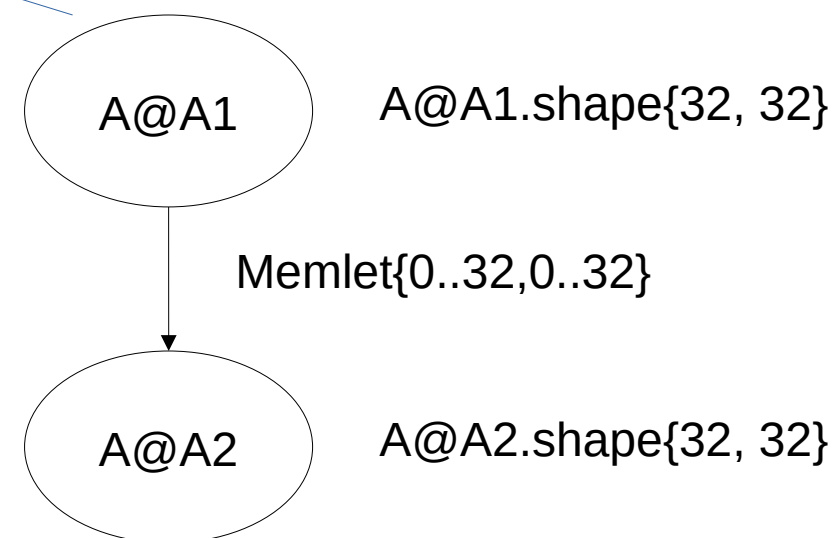
```
A_A1 = queue_A_A1.DeQueue<>()
```

```
For (...) {  
  int srcOffset = ..., int dstOffset = ...
```

```
  LoadDataParams;  
  LoadDataParams.repeatTimes = 32 / 16;  
  LoadDataParams.srcStride = 32 / 16;  
  LoadDataParams.ifTranspose = true;  
  
  LoadData(A_A2[dstOffset], A_A1[srcOffset],  
  LoadDataParams);  
  
}
```

```
queue_A_A2.Enqueue(A_A2);  
queue_A_A1.FreeTensor(A_A1);
```

1. Deque src



Anatomy of a Copy in SDFG:

```
A_A1 = queue_A_A1.DeQueue<>()
```

```
For (...) {  
  int srcOffset = ..., int dstOffset = ...
```

```
  LoadDataParams;  
  LoadDataParams.repeatTimes = 32 / 16;  
  LoadDataParams.srcStride = 32 / 16;  
  LoadDataParams.ifTranspose = true;
```

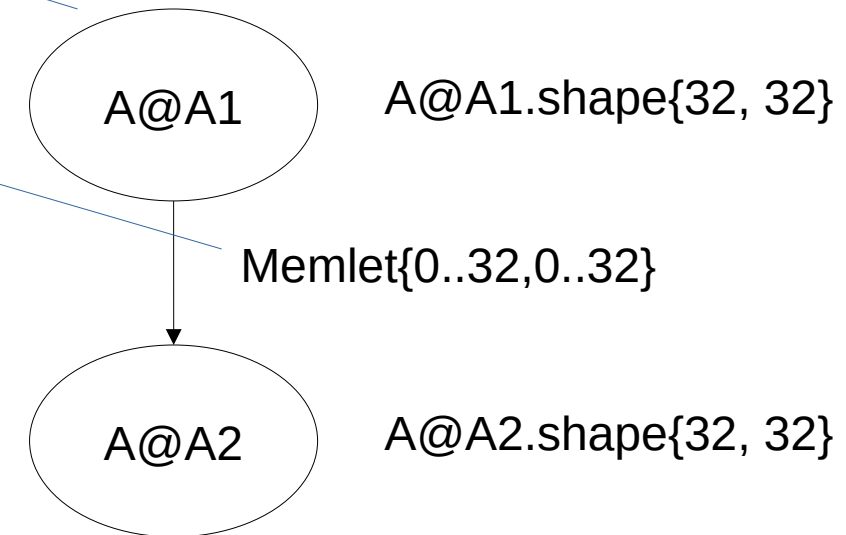
```
  LoadData(A_A2[dstOffset], A_A1[srcOffset],  
  LoadDataParams);
```

```
}
```

```
queue_A_A2.Enqueue(A_A2);  
queue_A_A1.FreeTensor(A_A1);
```

1. Dequeue src

2. generate the for loop to copy the
memlet shape as a unit of 16x16
blocks



Anatomy of a Copy in SDFG:

```
A_A1 = queue_A_A1.DeQueue<>()
```

```
For (...) {  
  int srcOffset = ..., int dstOffset = ...
```

```
  LoadDataParams;  
  LoadDataParams.repeatTimes = 32 / 16;  
  LoadDataParams.srcStride = 32 / 16;  
  LoadDataParams.ifTranspose = true;
```

```
  LoadData(A_A2[dstOffset], A_A1[srcOffset],  
  LoadDataParams);
```

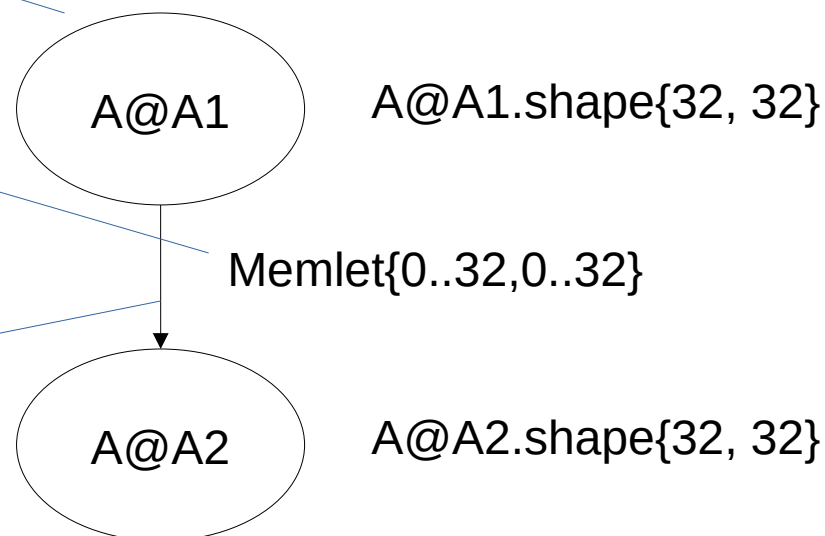
```
}
```

```
queue_A_A2.Enqueue(A_A2);  
queue_A_A1.FreeTensor(A_A1);
```

1. Dequeue src

2. generate the for loop to copy the memlet shape as a unit of 16x16 blocks

3. generate the copy call



Anatomy of a Copy in SDFG:

```
A_A1 = queue_A_A1.DeQueue<>()
```

1. Deque src

```
For (...) {  
  int srcOffset = ..., int dstOffset = ...
```

2. generate the for loop to copy the
memlet shape as a unit of 16x16
blocks

```
LoadDataParams;  
LoadDataParams.repeatTimes = 32 / 16;  
LoadDataParams.srcStride = 32 / 16;  
LoadDataParams.ifTranspose = true;
```

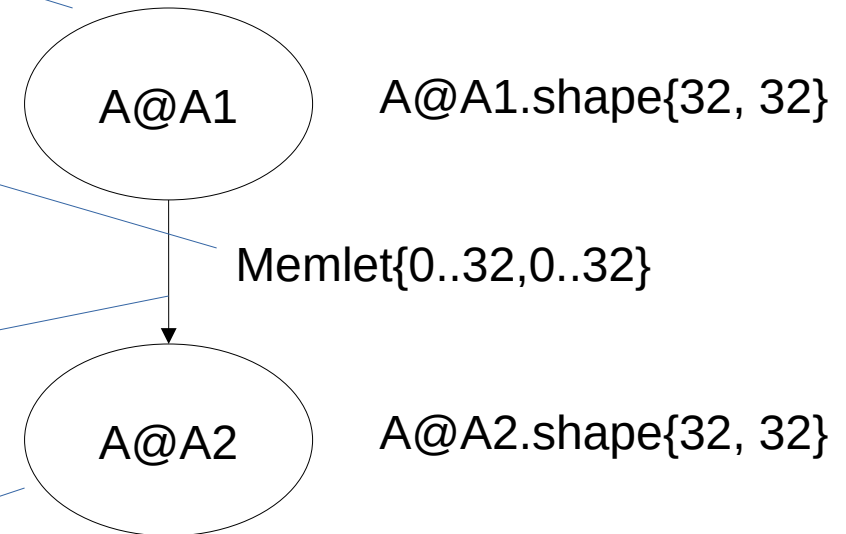
3. generate the copy call

```
LoadData(A_A2[dstOffset], A_A1[srcOffset],  
LoadDataParams);
```

```
}
```

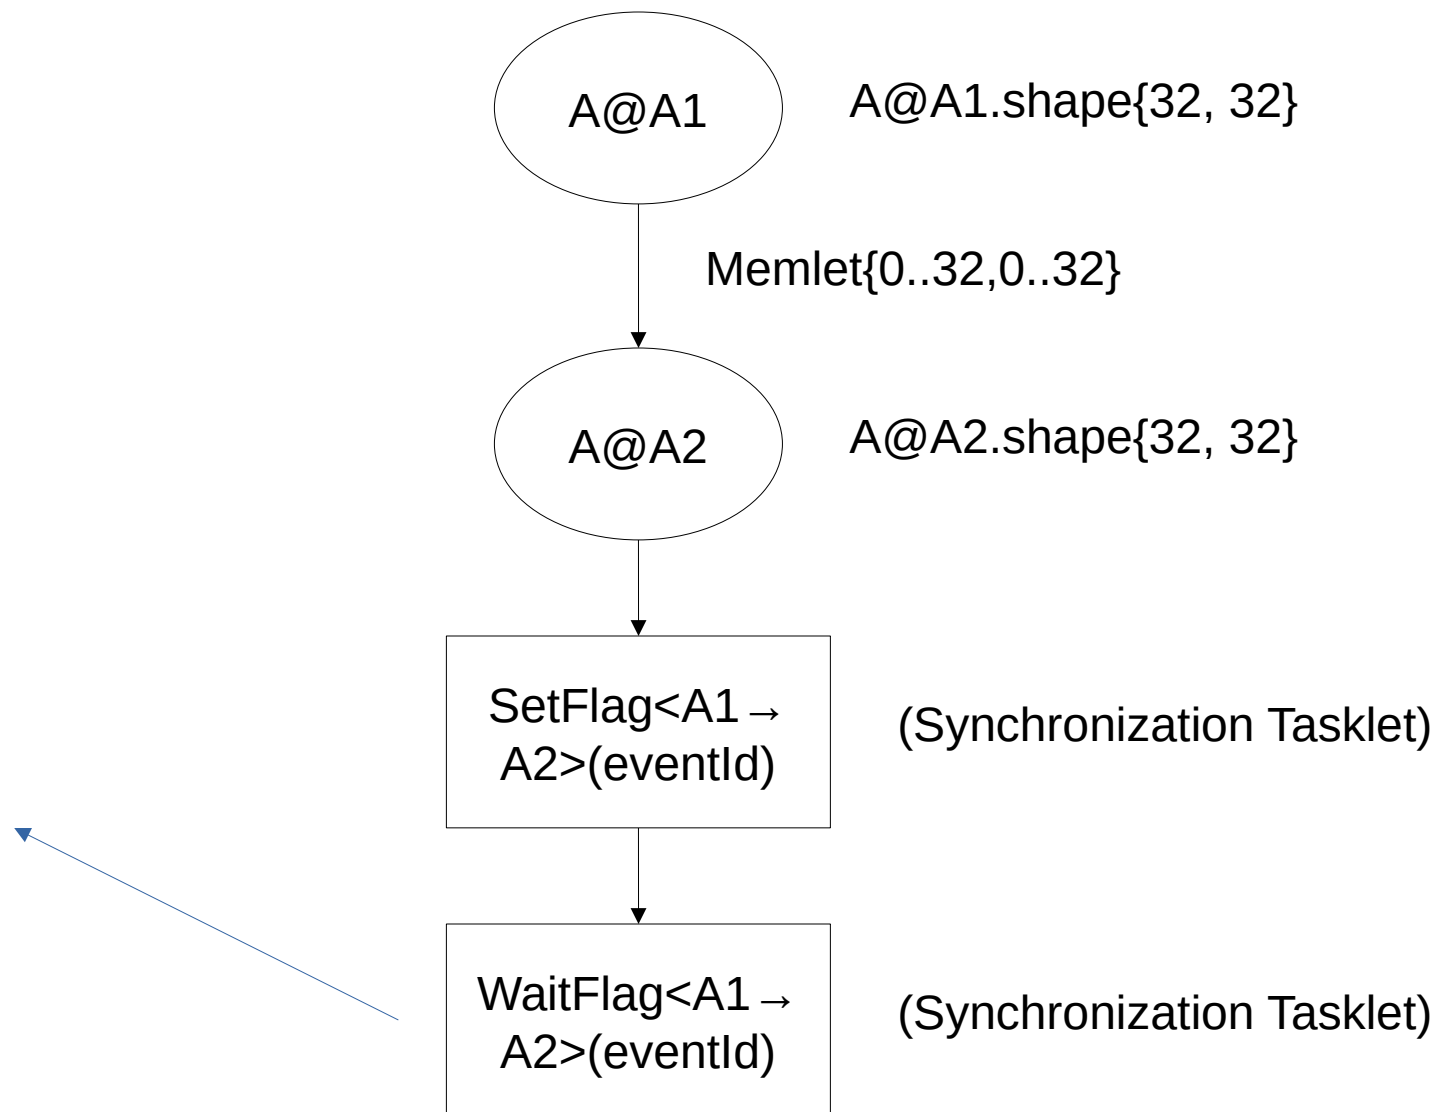
```
queue_A_A2.Enqueue(A_A2);  
queue_A_A1.FreeTensor(A_A1);
```

4. register copy, free src (alloc of
local tensor omitted)



Async Copy using AscendC Binary Semaphores:

Can be schedule to
be just before the
first read of A2.




Current Issues:

- Can generate syntactically correct AscendC code for the cube unit.
- Kernel timeouts (identified issues and currently working on fixes)

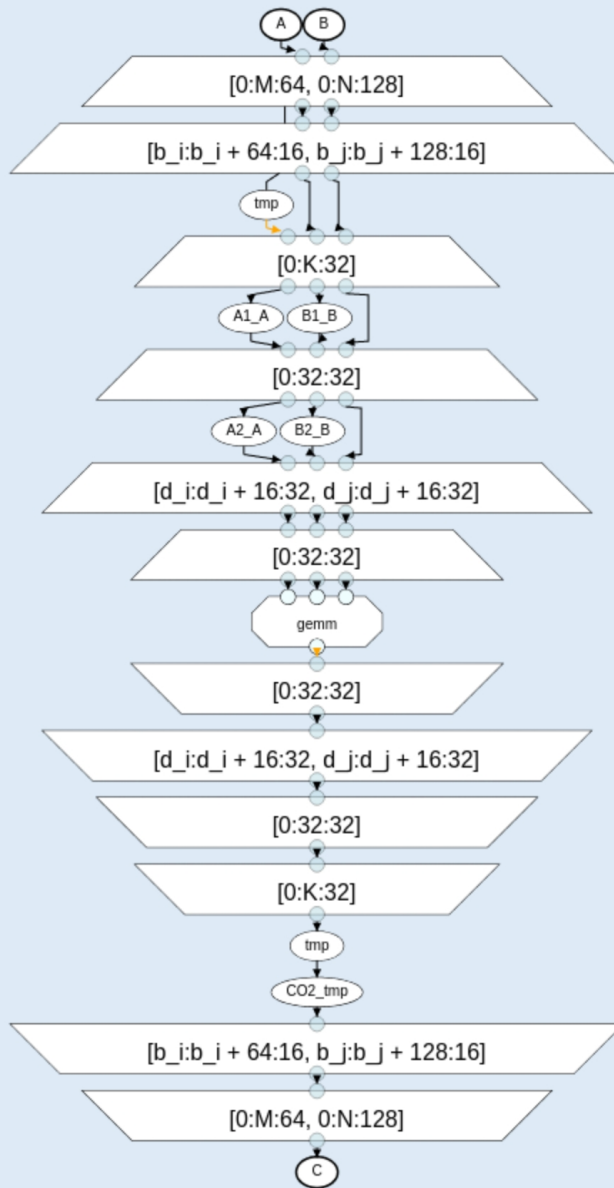
Current Issues:

$$C = A \times B$$



- Transformation (1) defines the per-thread computation domain ($T_M \times T_N$)
- Transformation (2) establishes the computation of the core-group (e.g., 32 for 910A, 20 for 910B4) domain ($B_M \times B_N$).
- The first tiling transformation (3) enables explicit data movement from global memory to A1 and B1 memory locations.
- The second tiling transformation (4) orchestrates movement from A1/B1 to A2/B2 respectively.

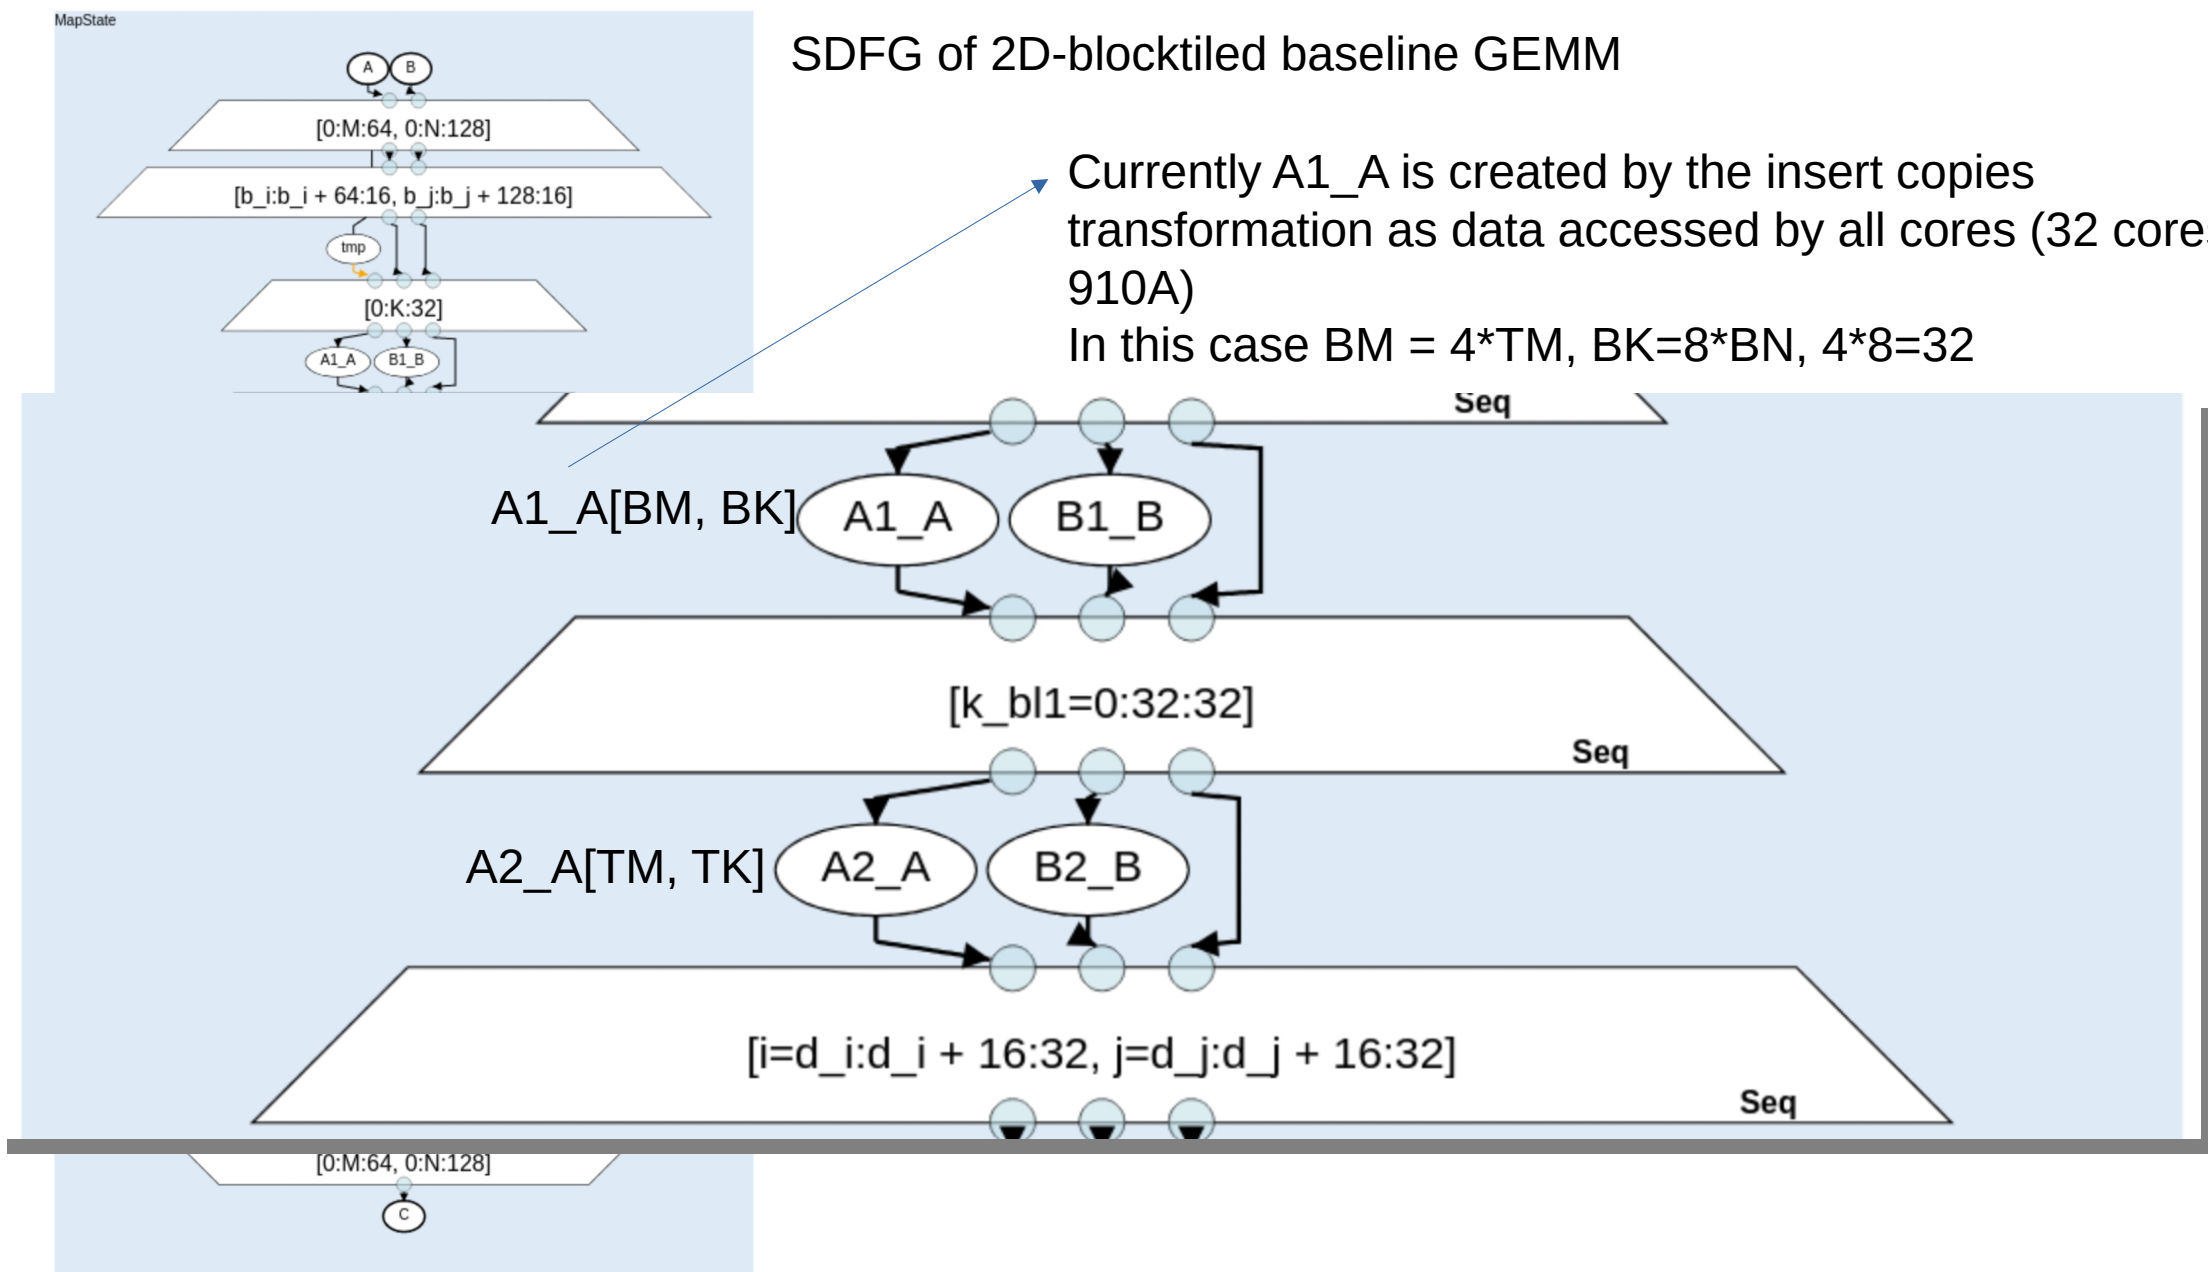
MapState



SDFG of 2D-blocktiled baseline GEMM

SDFG of 2D-blocktiled baseline GEMM

Currently A1_A is created by the insert copies transformation as data accessed by all cores (32 cores for 910A)
In this case $BM = 4 \cdot TM$, $BK = 8 \cdot BN$, $4 \cdot 8 = 32$



Current Issues (GitHub Issues Will Be Created Soon):

- For BlockTiling transformation (3: Global \rightarrow A1, Global \rightarrow B1) I created a hotfix to create the Glb \rightarrow A1 \rightarrow A2 and Glb \rightarrow B1 \rightarrow B2 directly within the same map (and only use thread-local tile) where A1 and A2 has the same size. As the current dimensions of A1 and B1 only make sense if cores can access other core's A1 or B1 storage.
 - If all cores copies the same memory locations from GM to A1 is there a broadcast mechanism?
 - Is it possible for a core to control/see the data another core loads to A1/B1 storage?
 - (The question is: if all levels of the memory locations expect GM are core-local, then will and how data locality effect the performance?)

Current Issues (GitHub Issues Will Be Created Soon):

Is moving partial deque from A1 to A2 allowed in the Que & Pipe syntax?

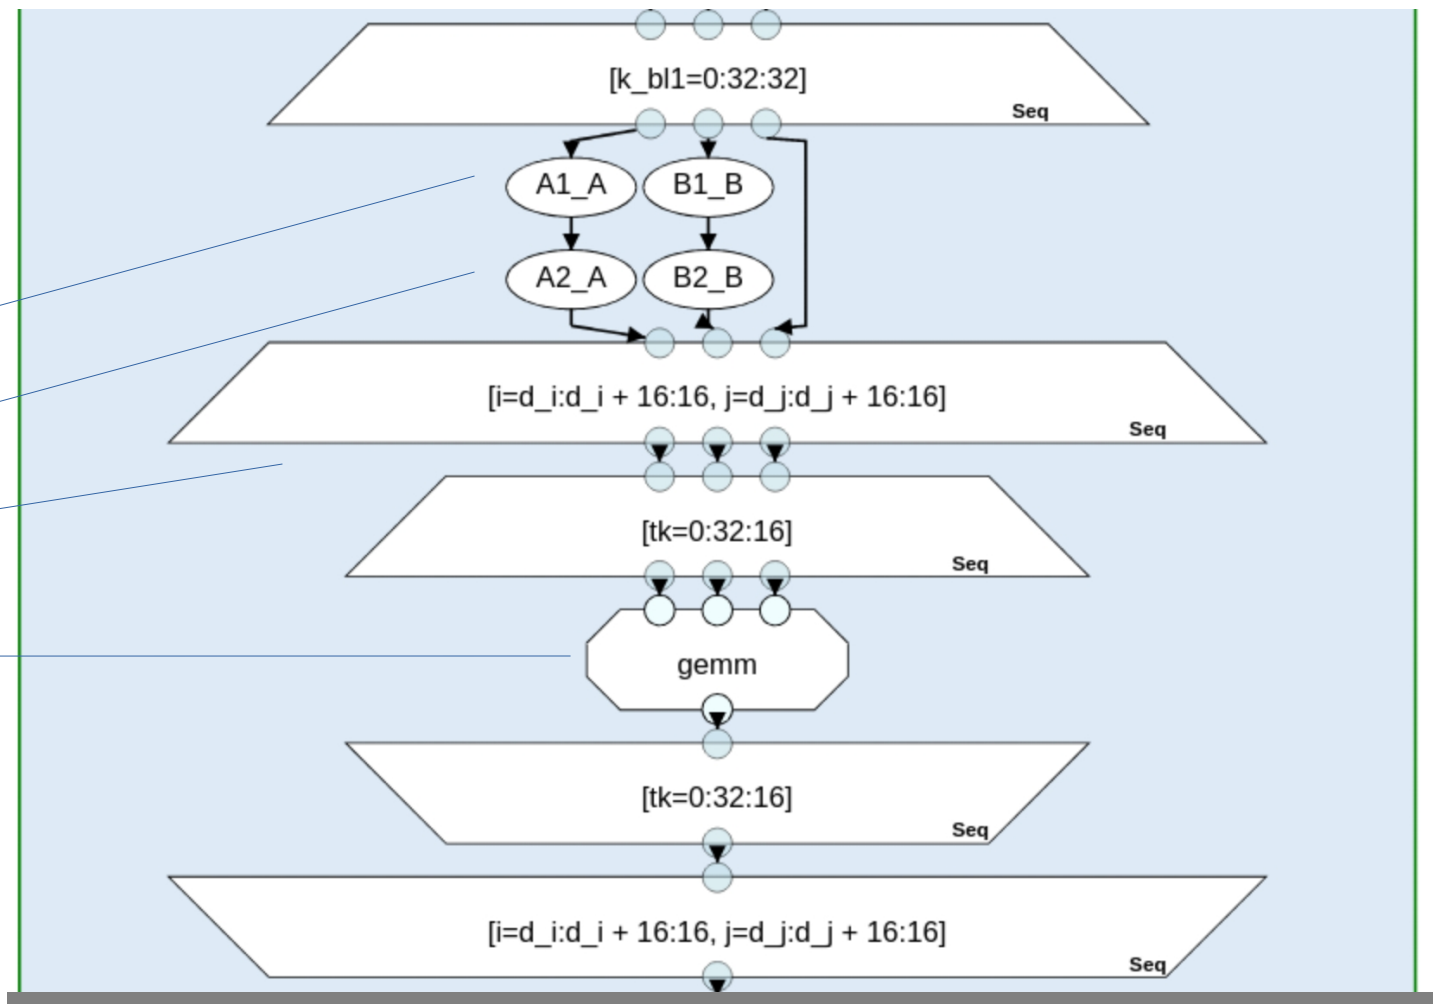
The Deque of A1_A and the following implicit synchronization indicates not.

What about if I use binary semaphores?

A1_A[16,32]

A2_A[16,32]

3 Calls the AscendC::Mmad
With M=N=K=16



Current Issues (GitHub Issues Will Be Created Soon):

- Then, do the steps $A1 \rightarrow A2$ and $B1 \rightarrow B2$ exist purely to make the layout conform with the layout required by the cube unit? (Such that the whole $A2$ data needs to be consumed within a call to the CubeUnit)
- Is it possible to load load the $A1$ to $A2$ in tiles within a for loop with data movement using Ques*

Current Issues (GitHub Issues Will Be Created Soon):

- From the DeQue and EnQueue operations I consider

```

1 pipe.InitBuffer(queue_A1_A, 1, 32 * 32 * sizeof(dace::float16));
2 pipe.InitBuffer(queue_A2_A, 1, 16 * 16 * sizeof(dace::float16));
3 ...
4 for (int i = 0; i < 4; ++i) { // Num Blocks
5     ...
6     A1_A = queue_A1_A.DeQueue<dace::float16>();
7     A2_A = queue_A2_A.AllocTensor<dace::float16>();
8     int xOffset = (i%2)*16;
9     int yOffset = (i/2)*16;
10
11     AscendC::LoadData2DParams A2_ALoadDataParams;
12     A2_ALoadDataParams.repeatTimes = 16 / 16; // 16 byze is the unit size
13     A2_ALoadDataParams.srcStride = 16 / 16;
14     A2_ALoadDataParams.ifTranspose = false; // No transpose for A
15
16     AscendC::LoadData(A2_A, A1_A[xOffset + yOffset * 16], A2_ALoadDataParams);
17     if (i == 3) {
18         queue_A2_A.Enqueue(A1_A);
19     } else {
20         queue_A2_A.FreeTensor(A1_A);
21     }
22     queue_A2_A.Enqueue(A2_A);
23     ...
24 }

```

This code would timeout / fail
due to out of bounds on
A2_A?

Next Steps:

- Complete GitHub Issues
- Fix the issues in the generated code
- If time permits:
- Work on the vectorized-stencil microbenchmark for the Vector Unit
- Introduce copies using binary semaphores for the vector unit